

Delivering easy-to-use frameworks to empower data-driven research

Dr. Rosa Filgueira, University of Edinburgh, EPCC



Background

- PhD Computer Science – Madrid/Carlos III
 - *Dynamic optimization techniques to enhance scalability and performance of MPI-based application*
- 5 years as a Postdoc – University of Edinburgh/DIR group
- 2 years as a Senior Data scientist at BGS

Currently – EPCC at the University of Edinburgh

- Data Architect activities across different domains/ Projects/
 - Scalability and performance of applications executed on HPC and Cloud resources
 - Scientific workflows, data-frameworks, containers and reproducibility tools, etc
 - Research activities

Introduction

Big Data Sciences Era, Data Intensive Computing applications



Scientific fields →
Data-driven

Astronomy, Geosciences,
Meteorology,, Bioinformatics



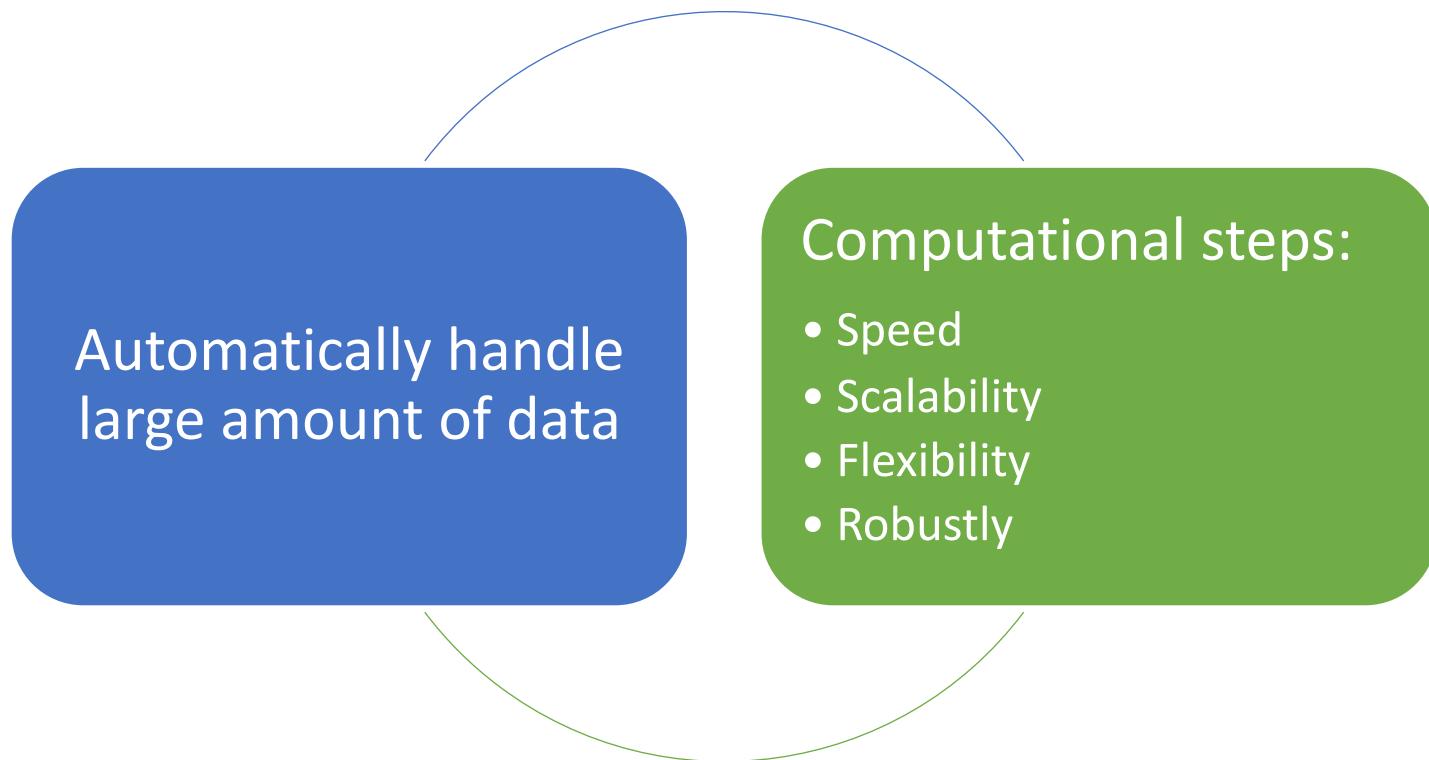
Common points

Big complex data sets
Need to be analysed
Numerous software tools
Data transformation and visualisation

Introduction

Big Data Sciences Era, Data Intensive Computing applications

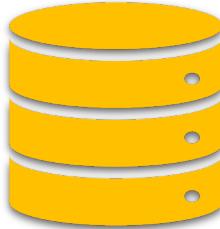
Data analytic frameworks and Computing environments



Delivering easy-to-use frameworks to empower data-driven research



Seismology



Scientific Workflows:

dispel4py

Pegasus

CWL



Computing environments:

HPC Clusters

Cloud

Why Scientific Workflows ?

General Features

Abstraction, scientists can focus on their research and not computation management

Easy composition and execution

Enables parallel, distributed computations

Different types

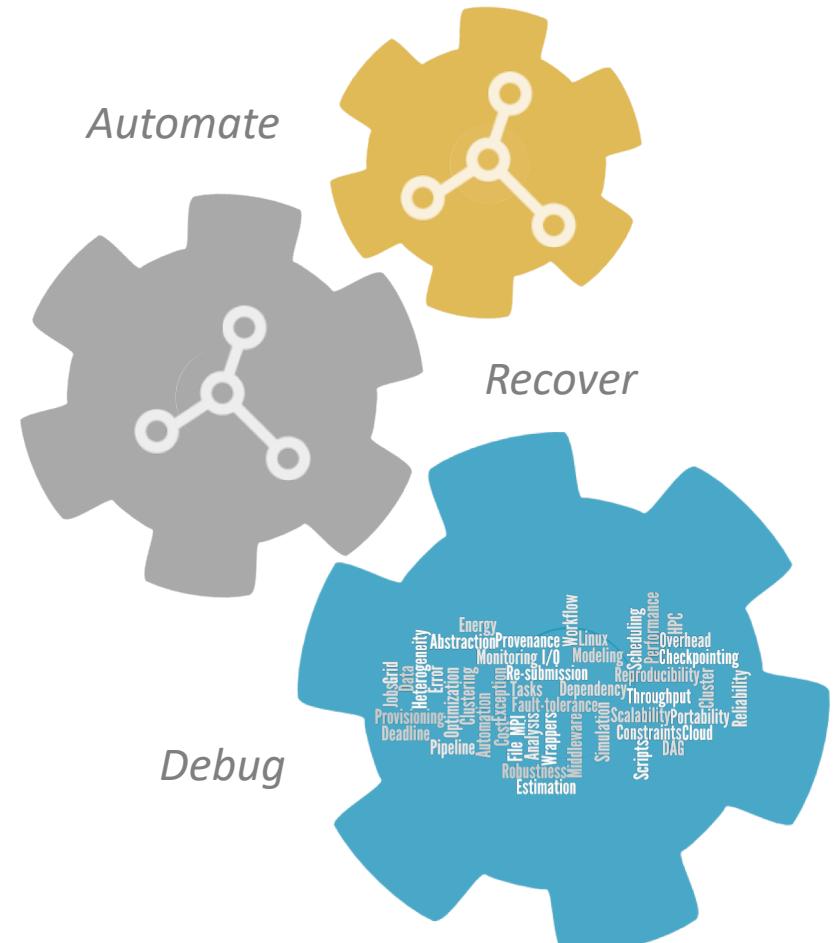
abstract vs. concrete

task-flow vs. data-flow

files vs. stream-based

Workflow Management Systems (WMS)

Provide tools to generate the scientific workflow



WORKFLOW MANAGEMENT SYSTEMS



Taverna

<https://taverna.incubator.apache.org>



KNIME

<https://www.knime.org>



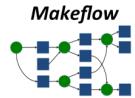
Kepler

<https://kepler-project.org>



VisTrails

<http://vistrails.org>



Makeflow

<http://ccl.cse.nd.edu/software/makeflow>



FireWorks

<https://pythonhosted.org/FireWorks>



dispel4py

<http://dispel4py.org>



Swift

<http://swift-lang.org>



Pegasus

<http://pegasus.isi.edu>



Nextflow

<https://www.nextflow.io>

Asterism Framework

Easy to understand, platform-independent, open-source

Simplifies the development applications running across multiple heterogeneous

How ?

Combining the strengths of



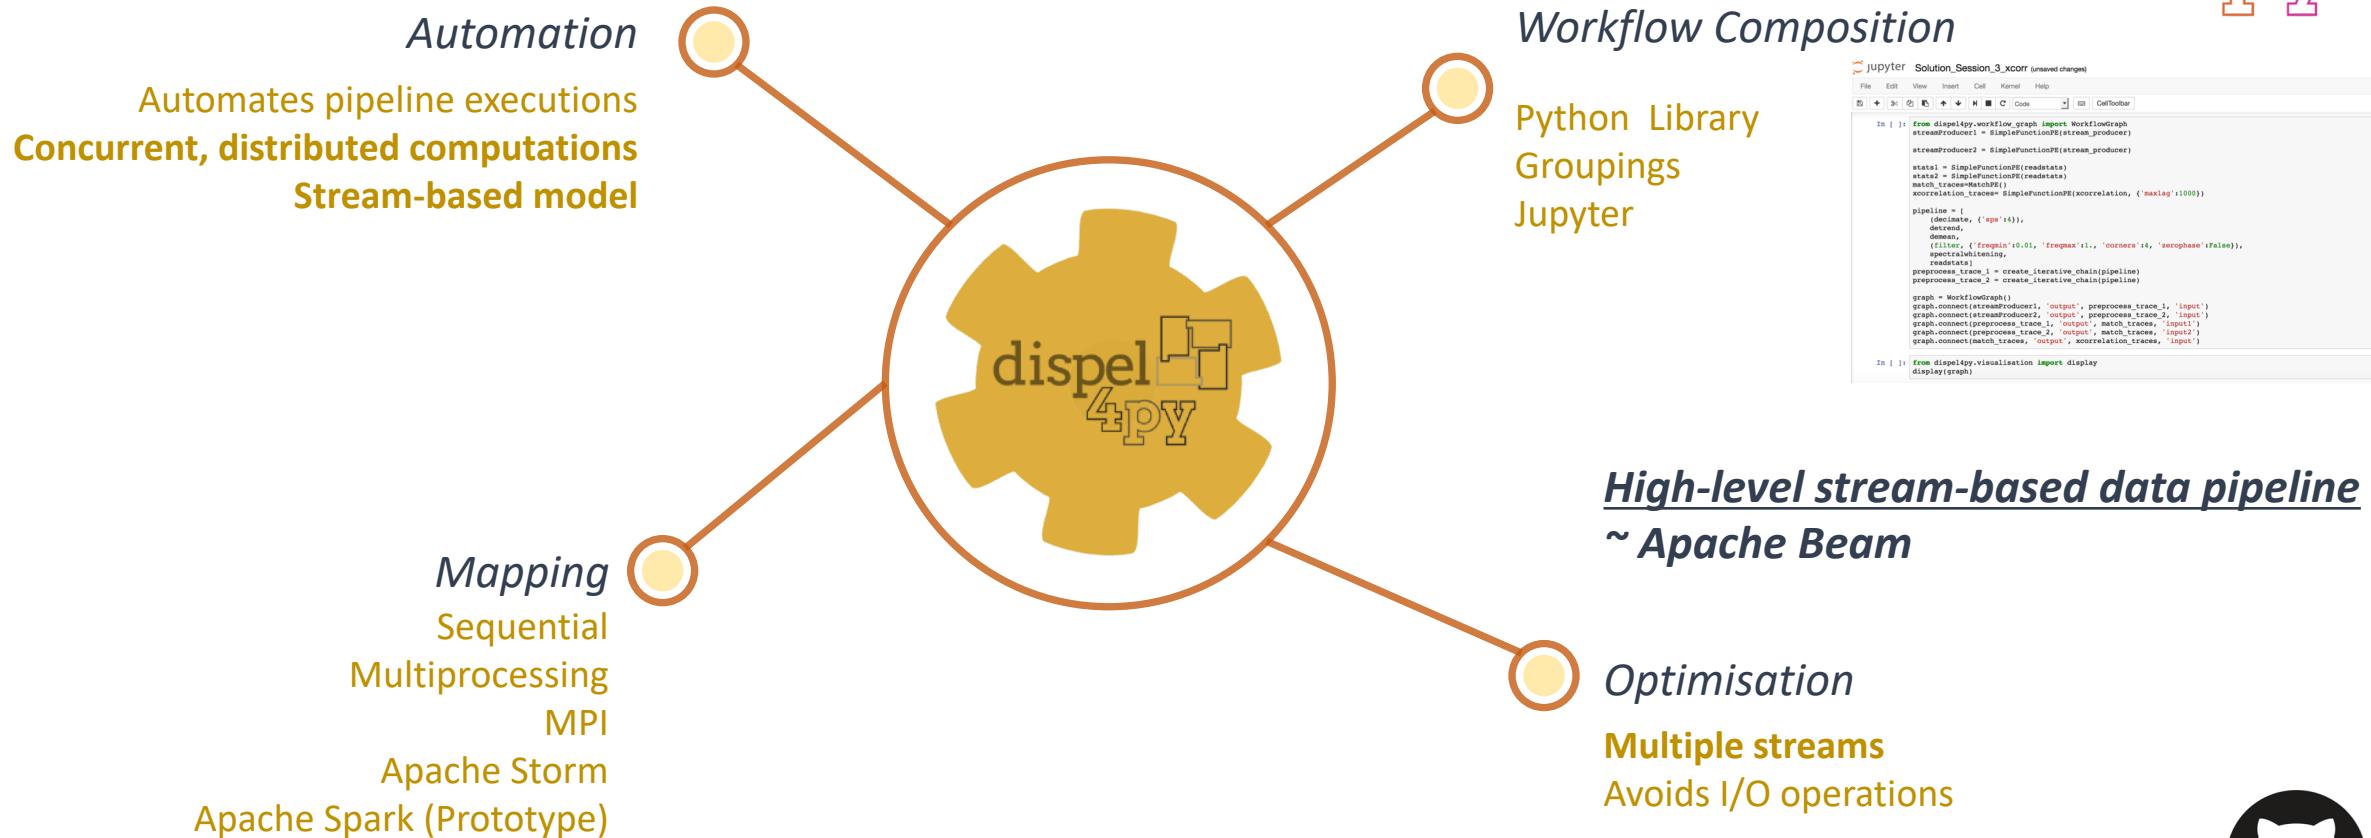
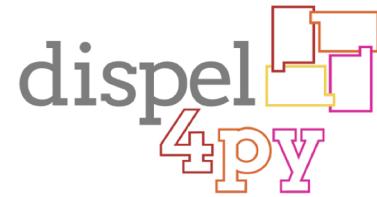
Traditional WMS
Pegasus

New stream-based data-flow systems
dispel4py

*Ewa Deelman, ISI-USC, US
Rafael Ferreira da Silva, ISI-USC, US*

*Rosa Filgueira, BGS-NERC, UK
Amrey Krause, EPCC-UoE, UK
Malcolm Atkinson, Informatics-UoE, UK*

dispel4py parallel stream-based dataflow system



A screenshot of a Jupyter Notebook interface. The code cell contains Python code for defining a workflow graph, connecting stream producers, and creating iterative chain pipelines. The output cell shows the resulting graph visualization.

```
In [1]: from dispel4py.workflow_graph import WorkflowGraph
streamProducer1 = SimpleFunctionP(stream_producer)
streamProducer2 = SimpleFunctionP(stream_producer)
readData1 = SimpleFunctionP(readData)
stats1 = SimpleFunctionP(readData)
xcorr = SimpleFunctionP(xcorrelation, {'maxlag':1000})
pipeline = [
    (decimate, {'npp':14}),
    (detrend, {}),
    (filter, {'fregmin':10.0, 'fregmax':11., 'corner':4, 'zerophase':False}),
    (spectralEstimating, {}),
    (readData),
    (process_trace_1 = create_iterative_chainpipeline),
    (process_trace_2 = create_iterative_chainpipeline),
    graph = WorkflowGraph()
]
graph.connect(streamProducer1, 'output', preprocess_trace_1, 'input')
graph.connect(streamProducer2, 'output', preprocess_trace_2, 'input')
graph.connect(preprocess_trace_1, 'output', match_traces, 'input1')
graph.connect(preprocess_trace_2, 'output', match_traces, 'input2')
graph.connect(match_traces, 'output', xcorr, 'input')
In [2]: from dispel4py.visualization import display
display(graph)
```



Key-features: Automatic parallelization/mappings, concurrent & stream-based



dispel4py parallel stream-based dataflow system

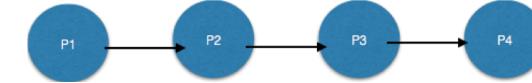


Graph

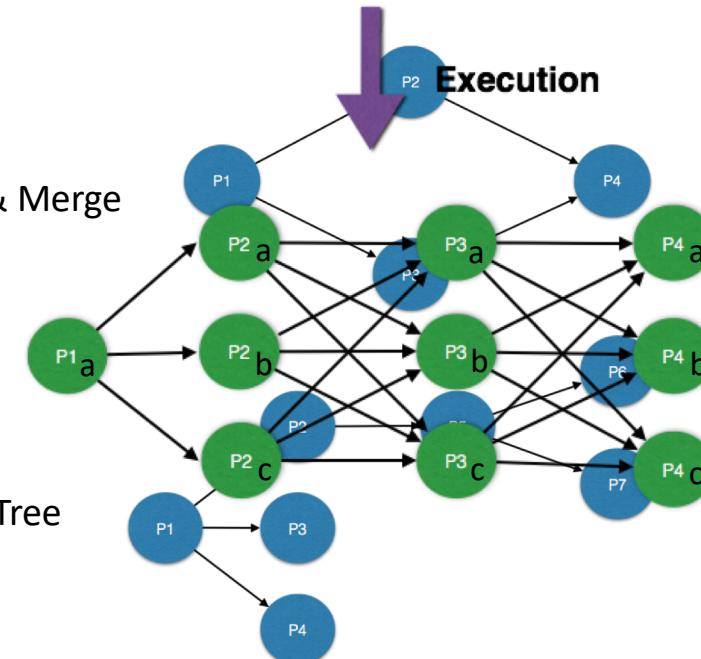
- Connections among PEs
- Abstract workflow

+ Example of graphs

Pipeline



Split & Merge



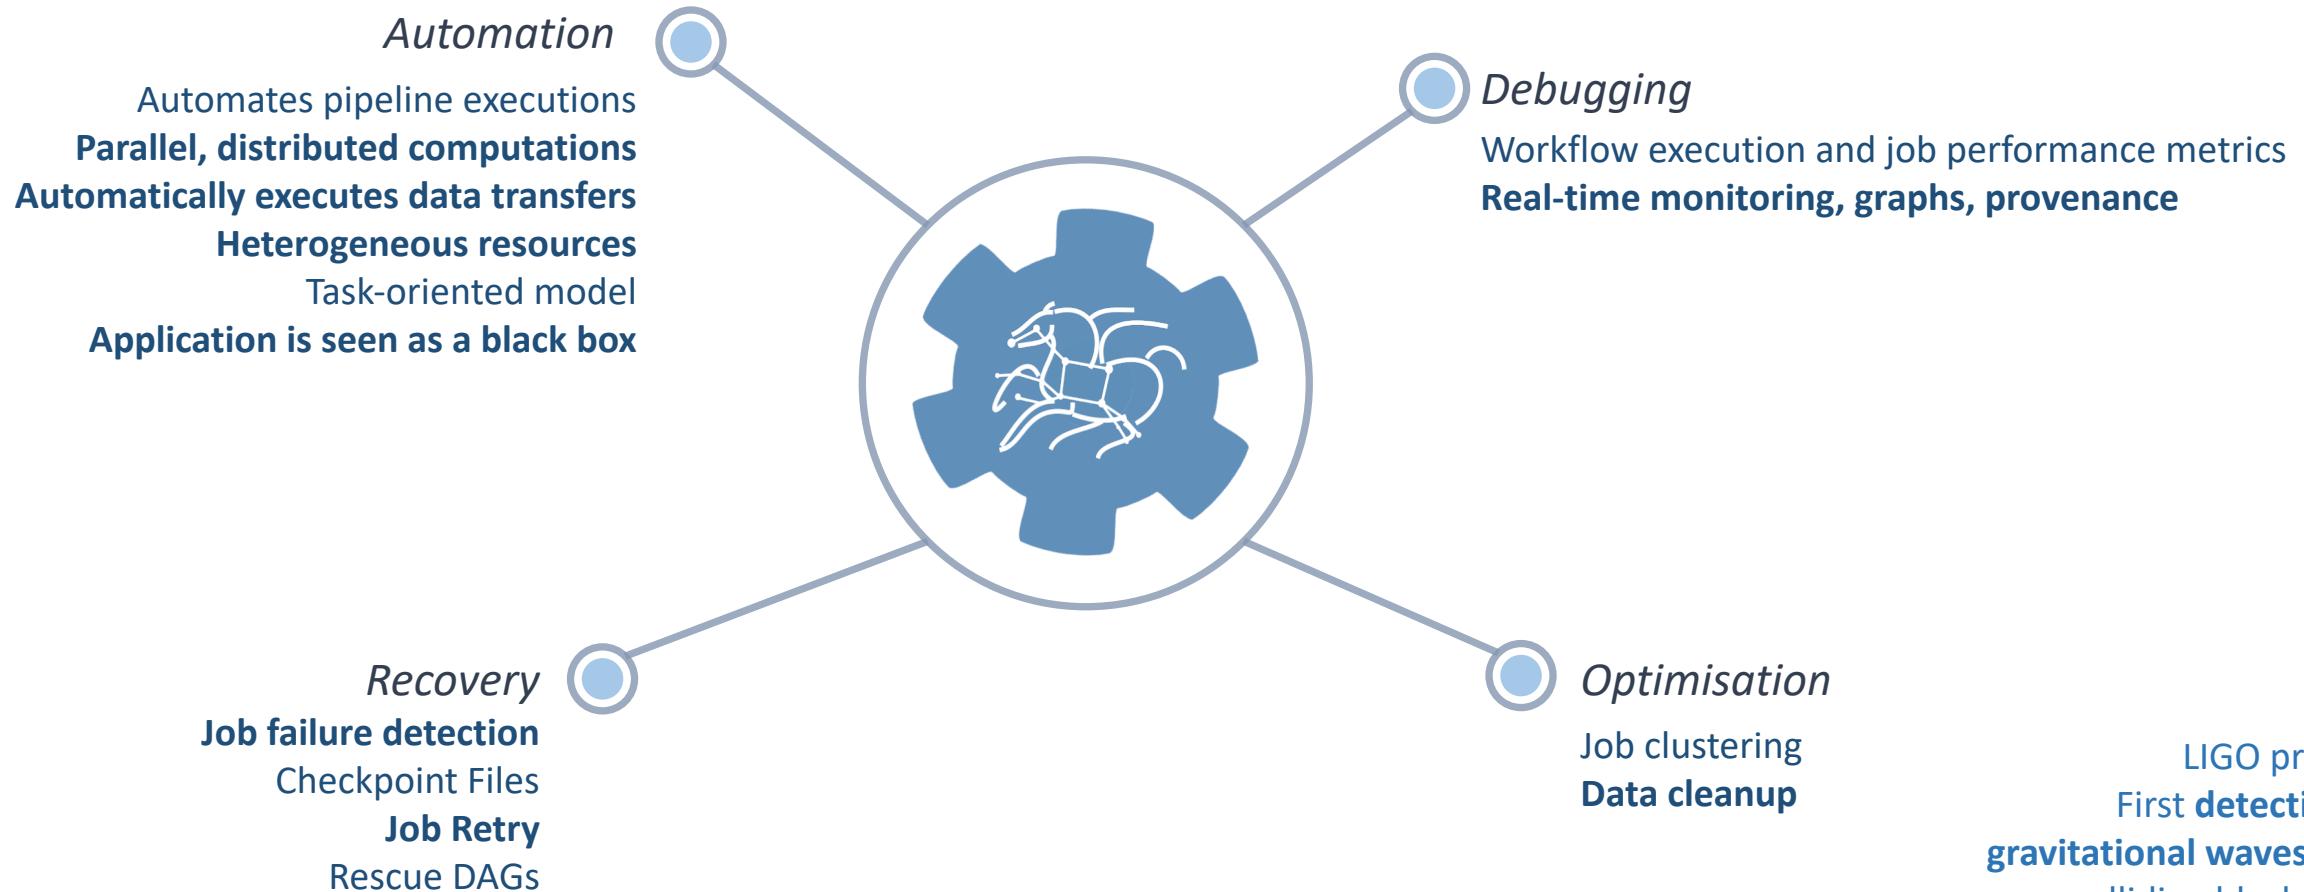
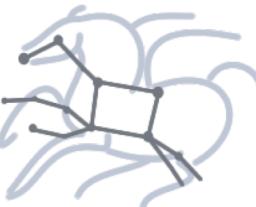
Tree

4 PEs & 10 processes



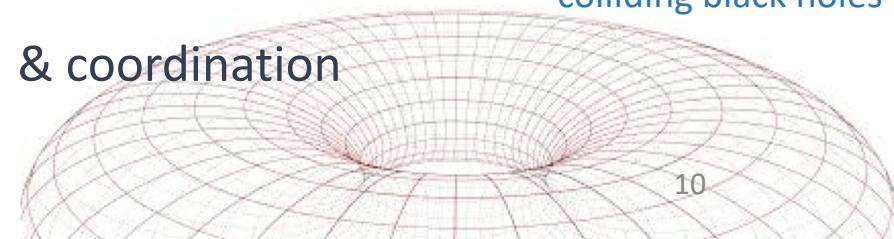
epcc

Pegasus workflow system

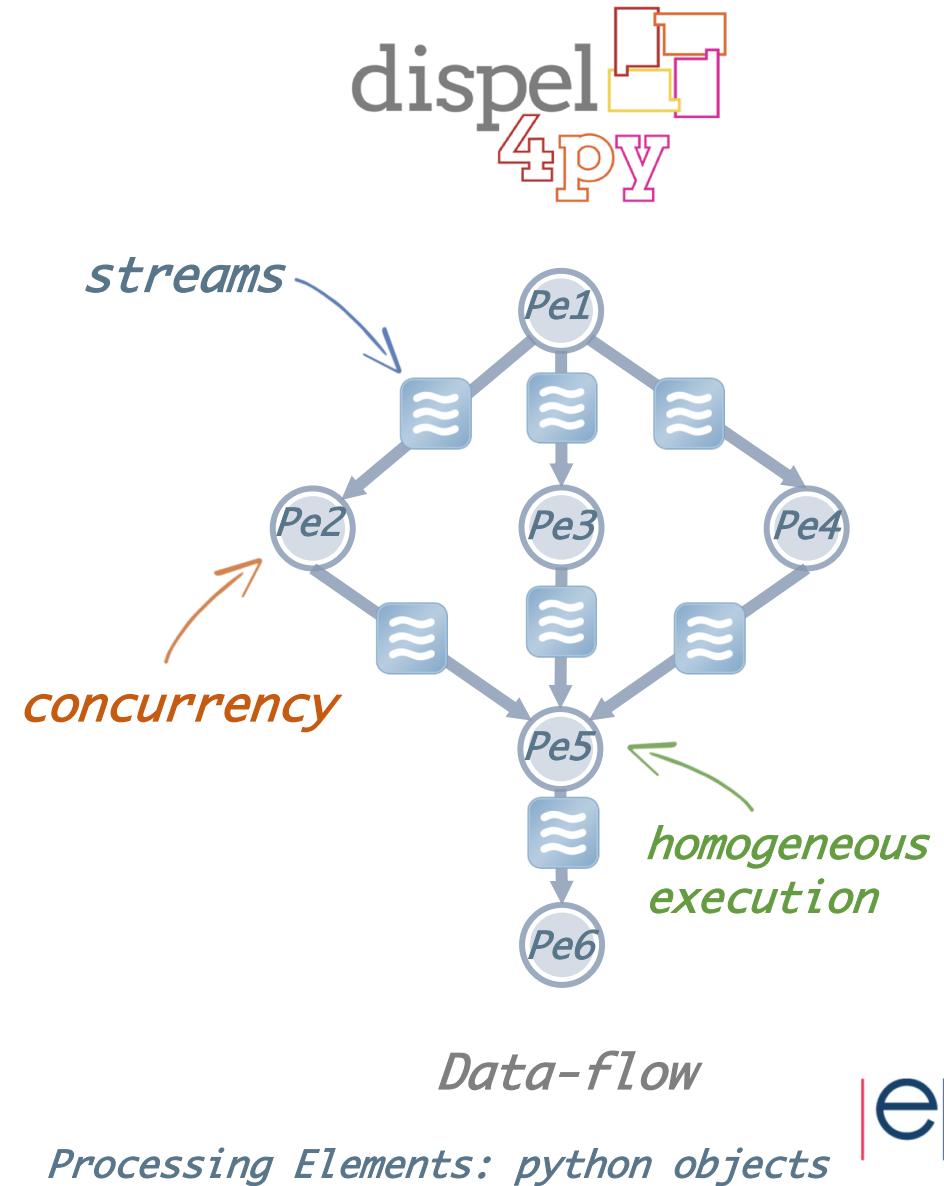
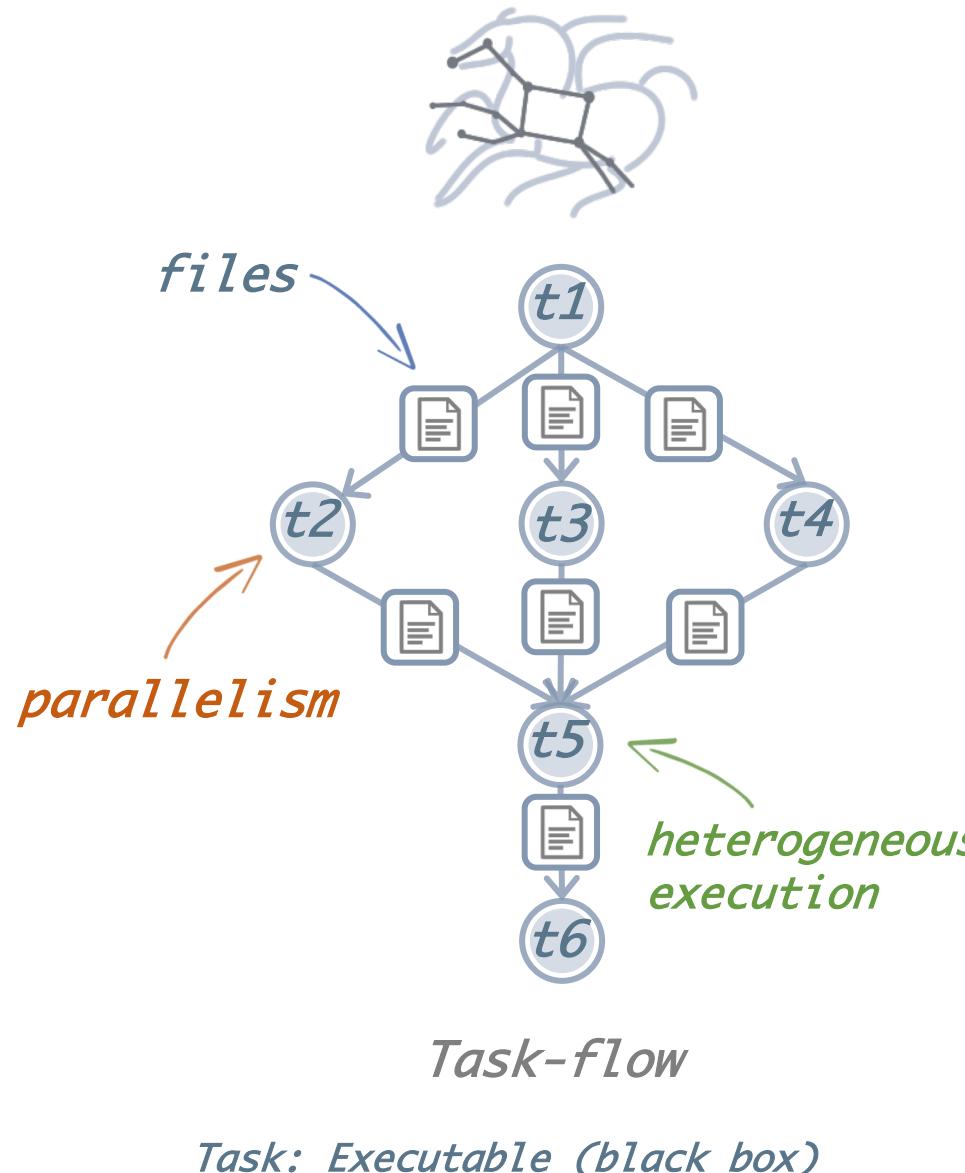


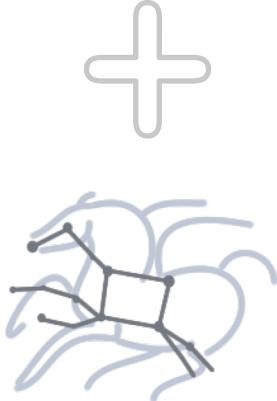
Key-features: Automatic data movement, cleanup, heterogeneous & coordination

LIGO project:
First detection of
gravitational waves from
colliding black holes

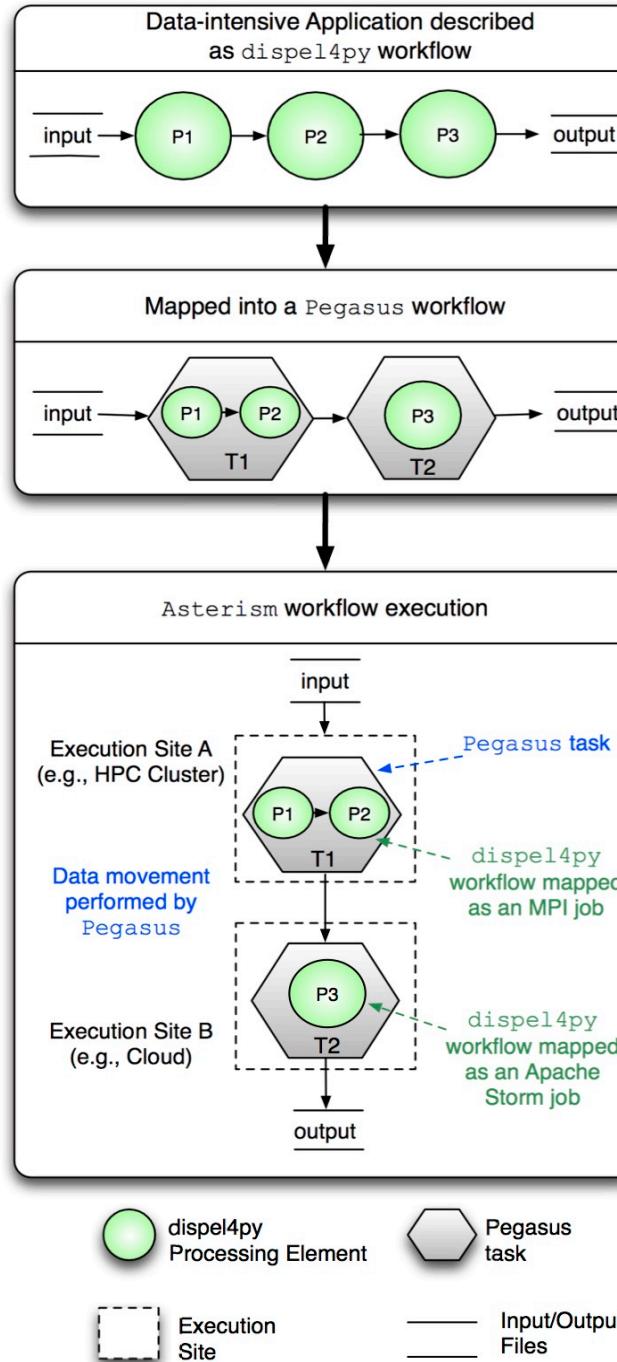


Complementary systems





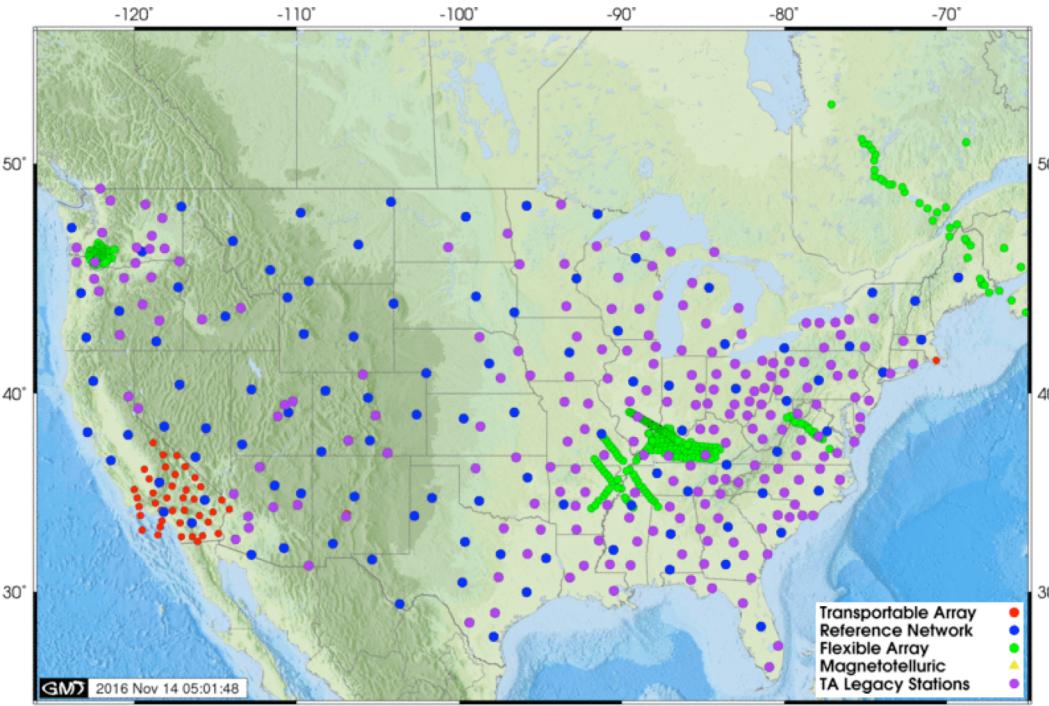
ASTERISM



dispel4py to represent different parts of applications

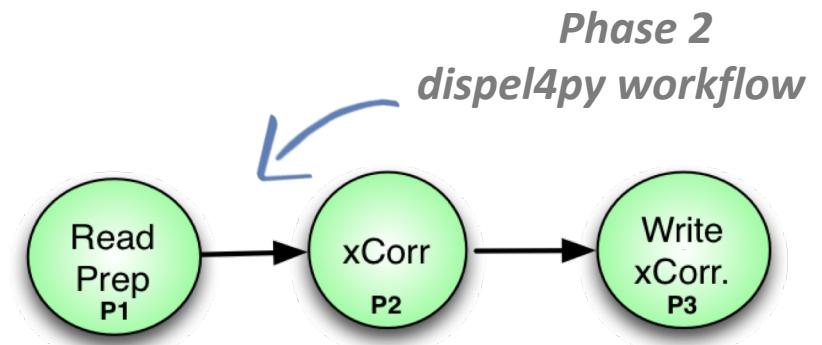
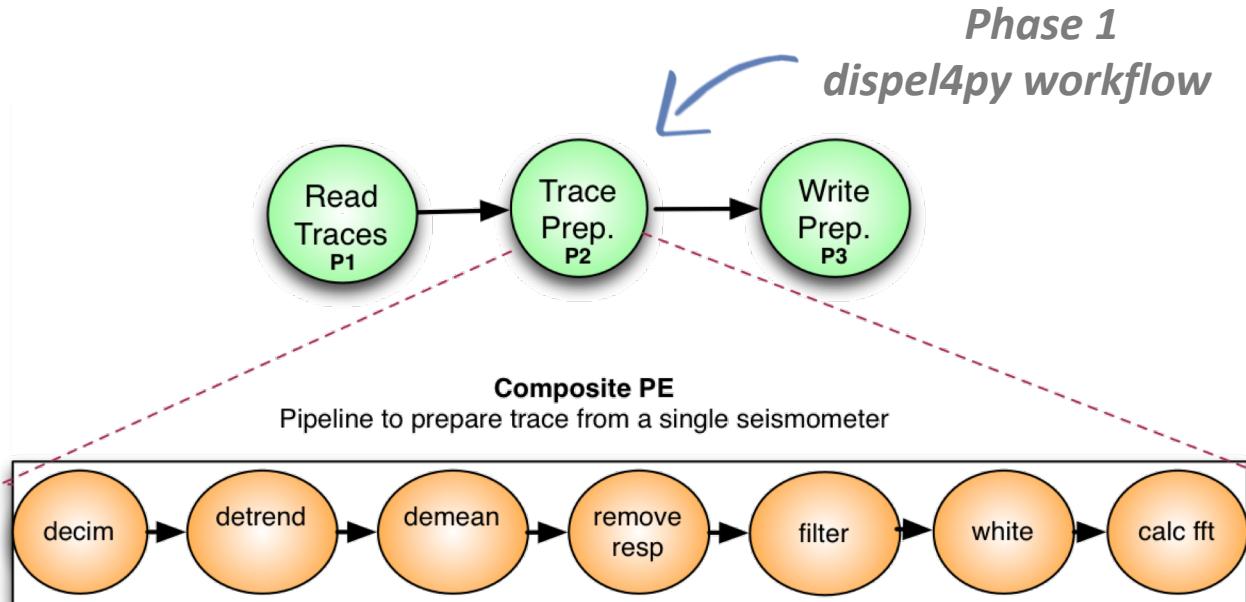
Pegasus to distribute and execute each dispel4py workflow

Seismic Ambient Noise Cross-Correlation



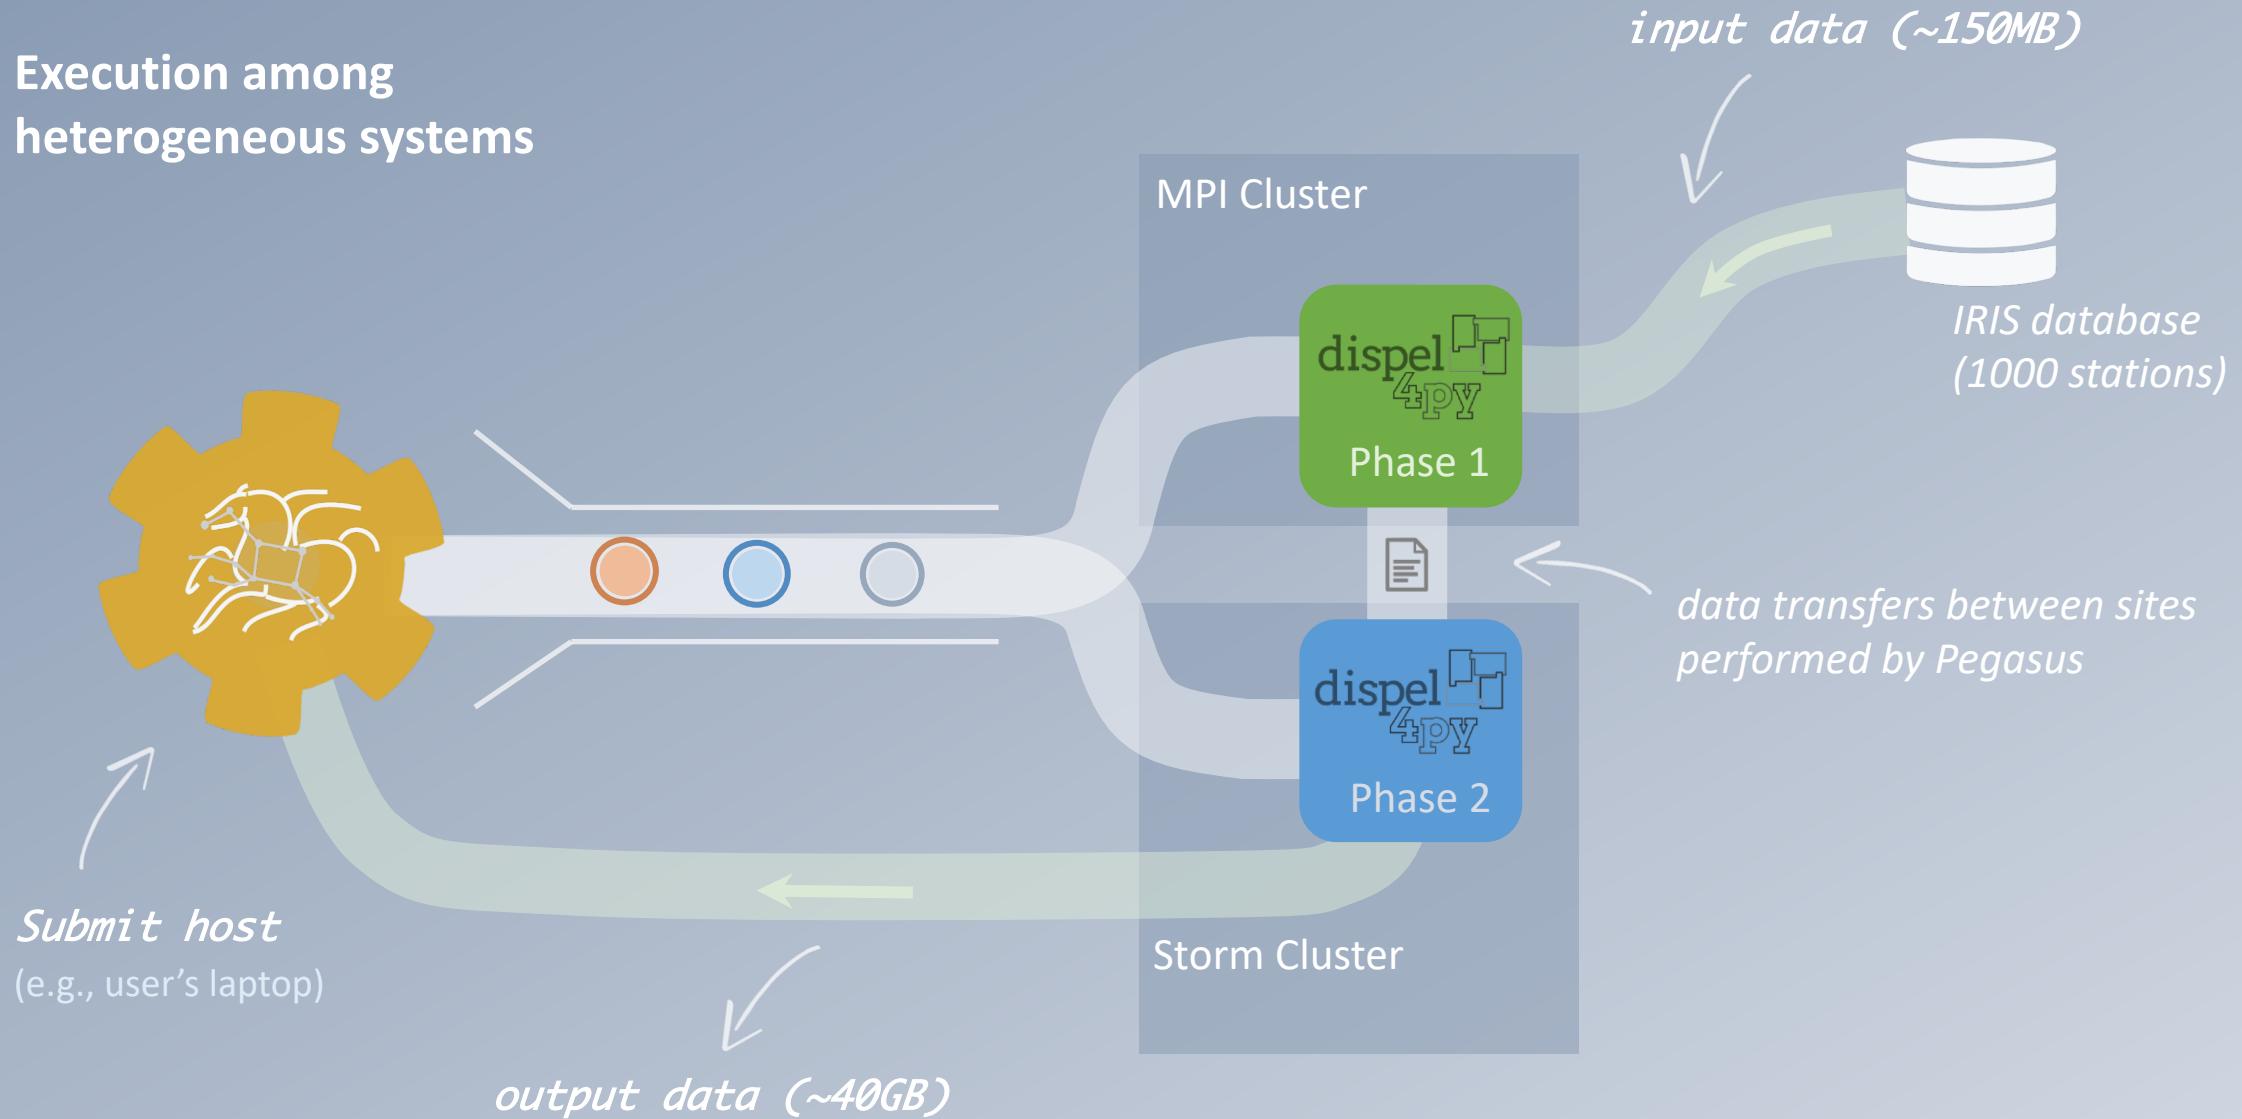
Preprocesses (Phase 1) and cross-correlates traces (Phase 2) from multiple seismic stations

VERCE project both phases on the same HPC resource



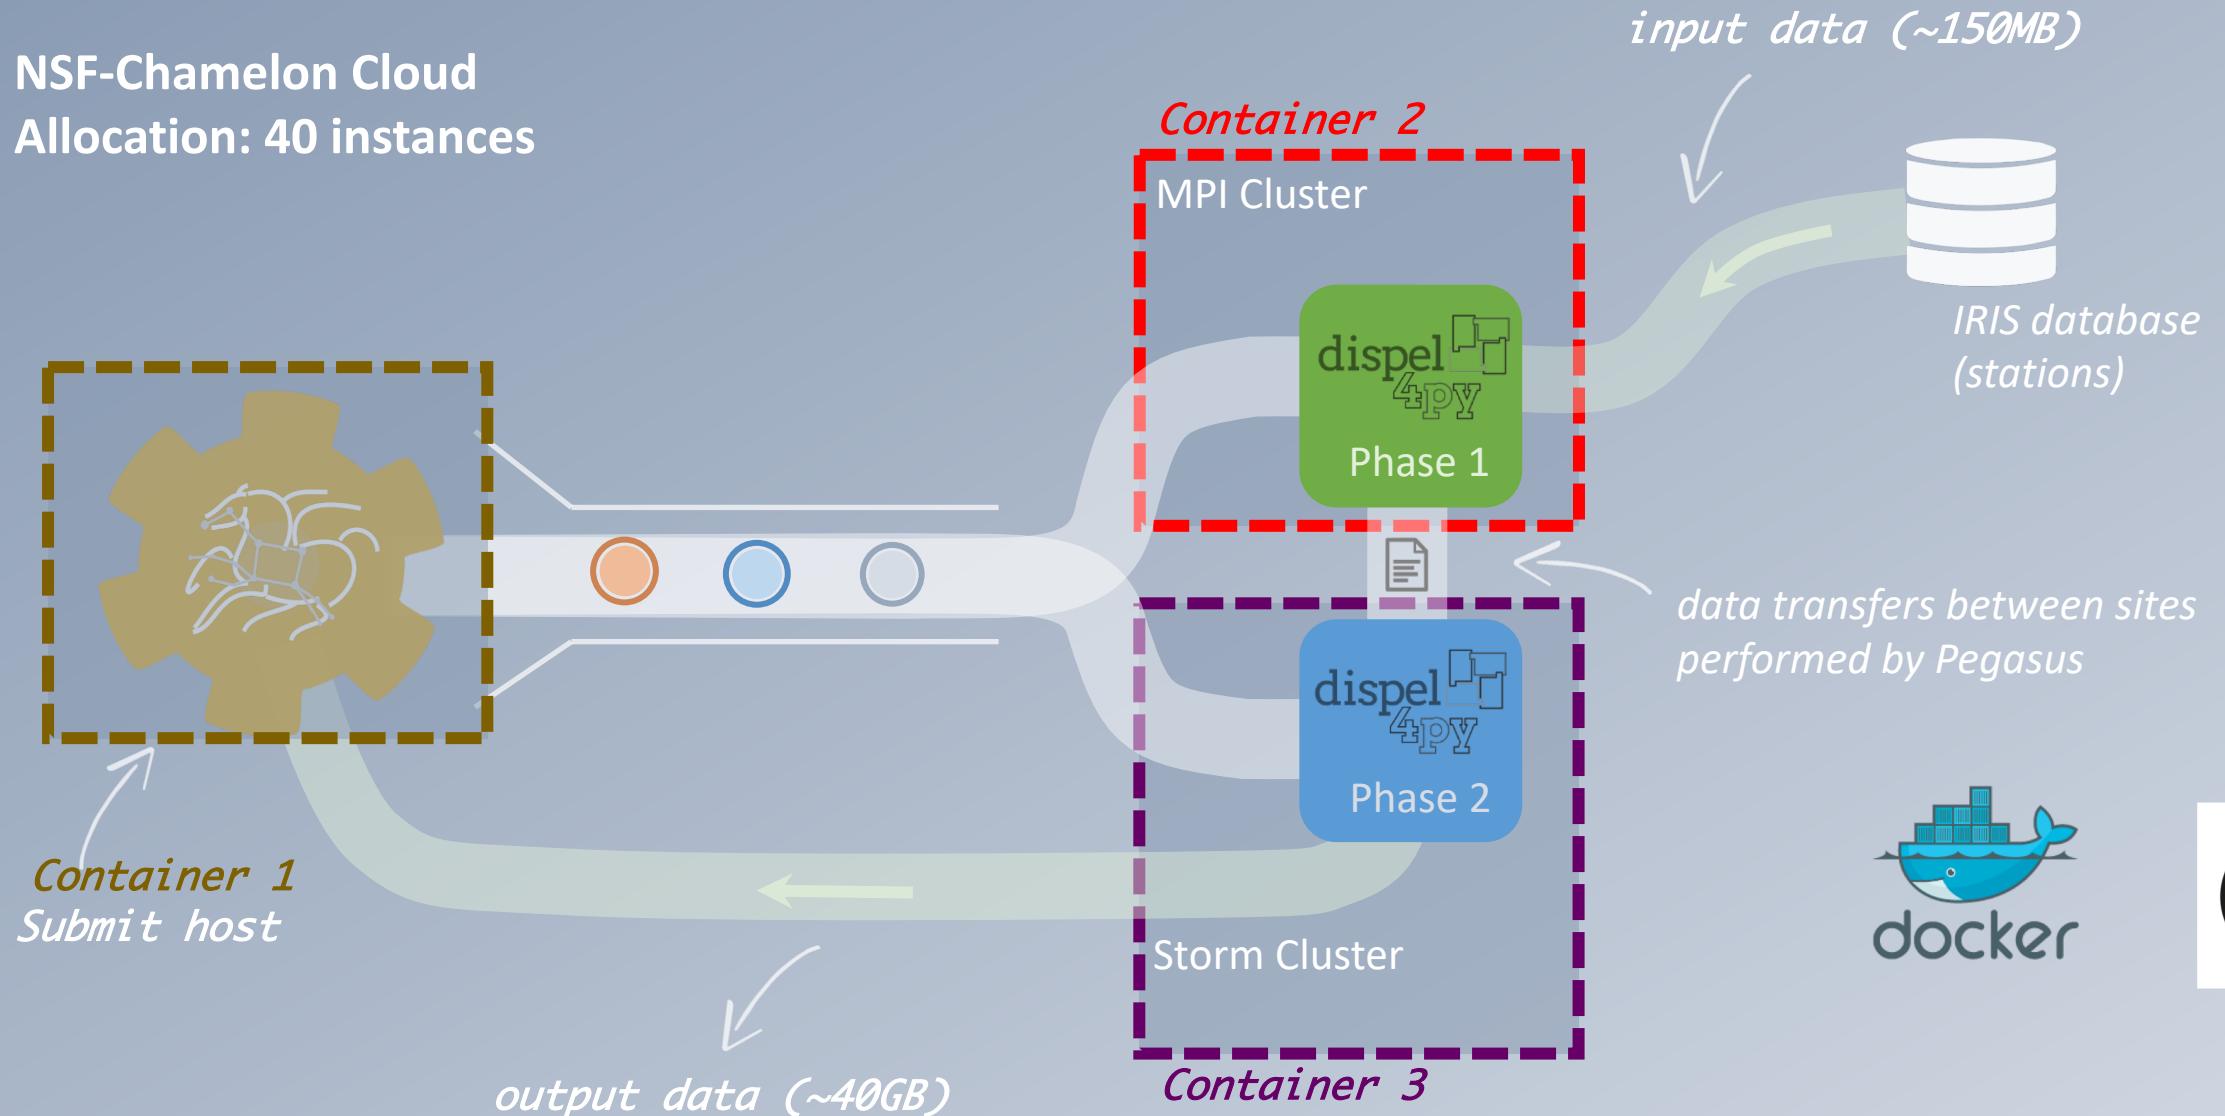
Evaluation- Seismic Ambient Noise Cross-Correlation

Execution among
heterogeneous systems



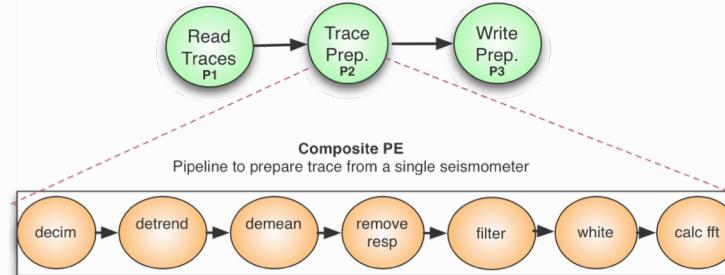
Evaluation- Seismic Ambient Noise Cross-Correlation

NSF-Chameleon Cloud
Allocation: 40 instances

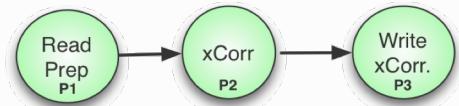


Reminder

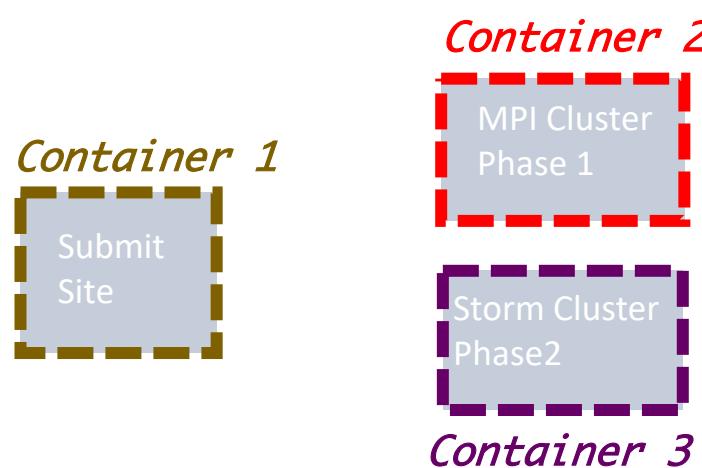
Data-Intensive Application



dispel4py preproc. (Phase 1)



dispel4py proc. (Phase 2)



NSF-Chameleon cloud
1 node – 40 cores

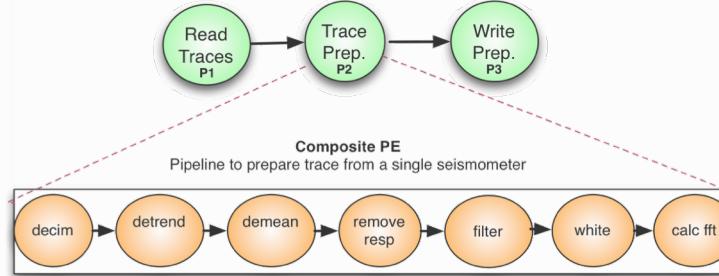


Dockerfiles to configure containers **images** → Stored in our **GitHub** → Linked to **DockerHub** → stored/share/download images

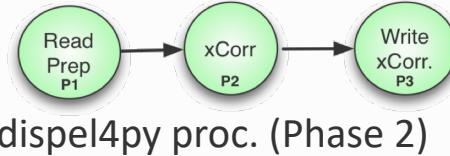


Reminder

Data-Intensive Application



dispel4py preproc. (Phase 1)



dispel4py proc. (Phase 2)

Container 2



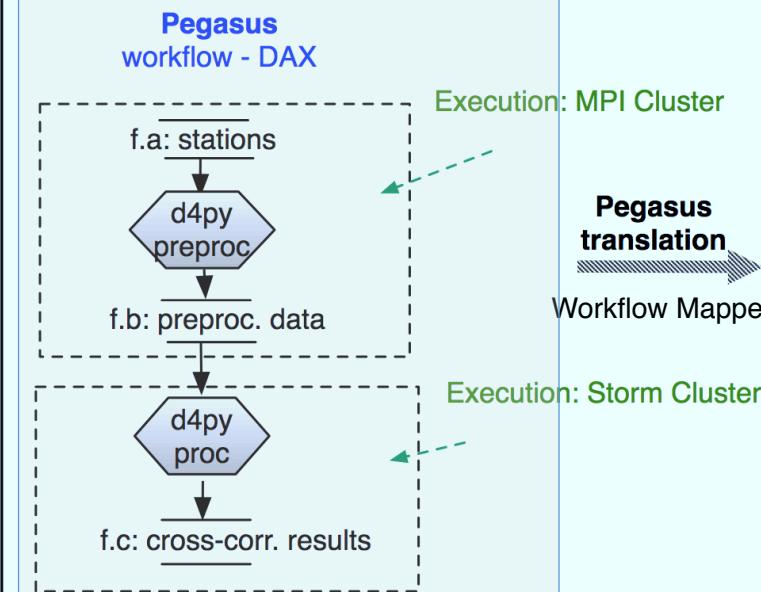
Container 1



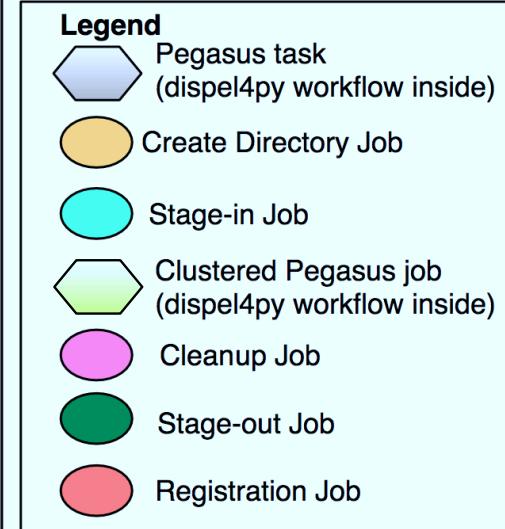
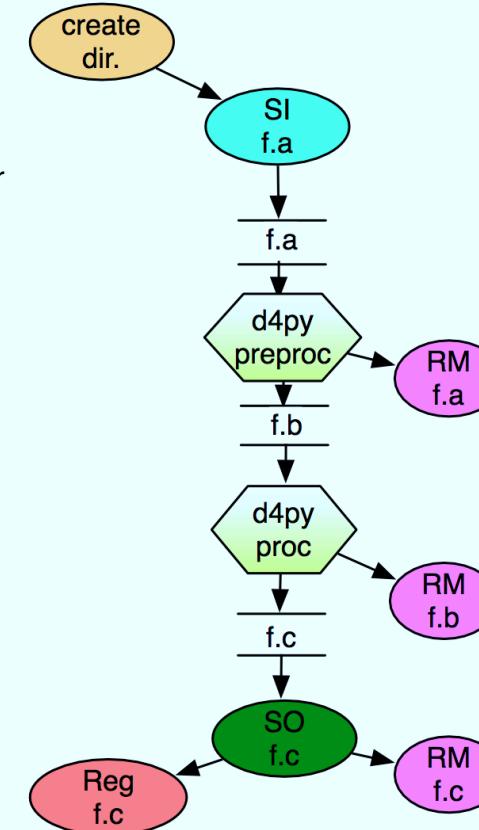
Container 3

Asterism Seismic Cross Correlation workflow

Execution environment -- Container 1: Pegasus, HTCondor, dispel4py



Pegasus executable workflow



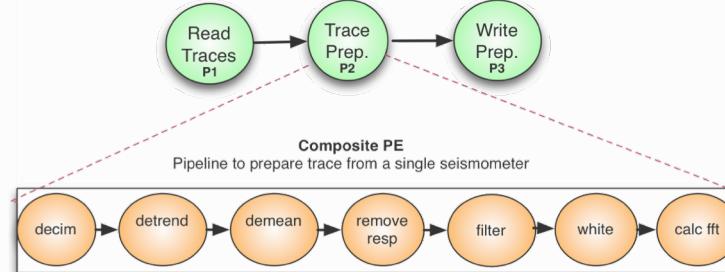
1 instance as Container 1

NSF-Chameleon cloud
1 node – 40 cores

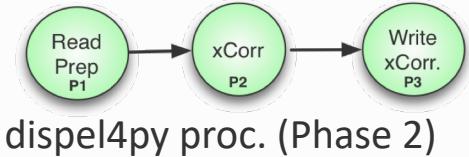


Reminder

Data-Intensive Application



dispel4py preproc. (Phase 1)



dispel4py proc. (Phase 2)

Container 2



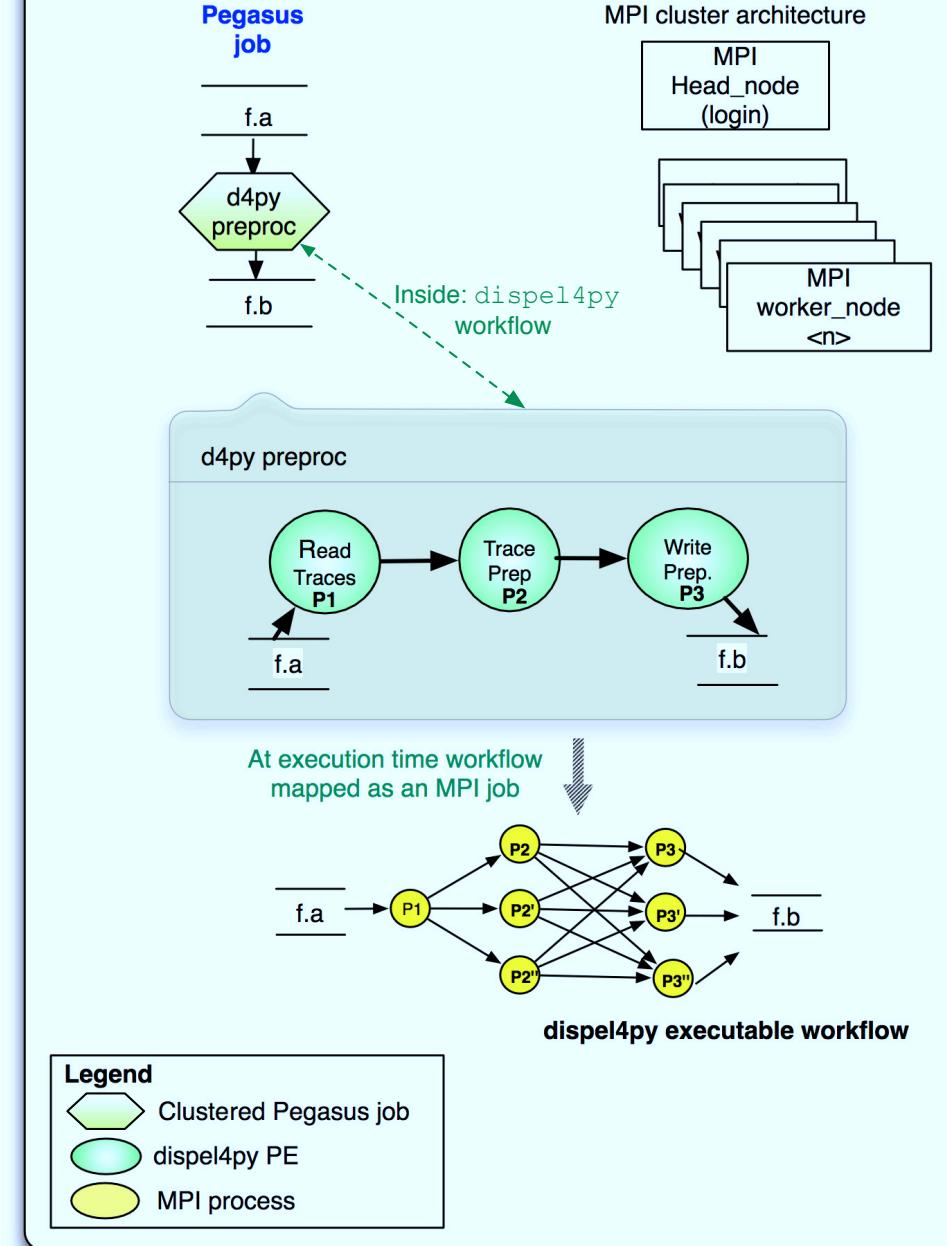
Container 1

Submit Asterism



Container 3

Execution environment -- Container 2- MPI cluster, dispel4py, Obspy



**1 instance as
Container 2
(MPI head node)**

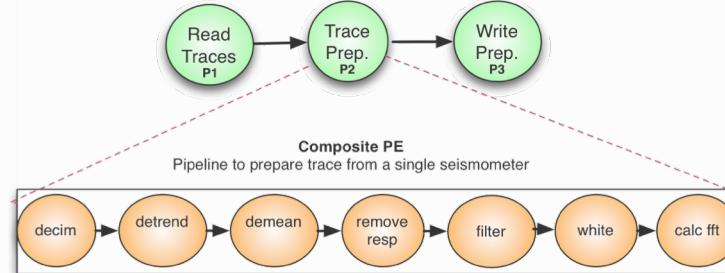
**16 instances as
Container 2
(MPI workers)**

NSF-
Chameleon
cloud
1 node –
40 cores

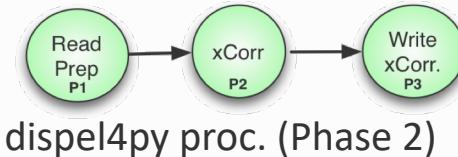


Reminder

Data-Intensive Application



dispel4py preproc. (Phase 1)



dispel4py proc. (Phase 2)

Container 2

MPI Cluster

Phase 1

Storm Cluster

Phase 2

Container 3

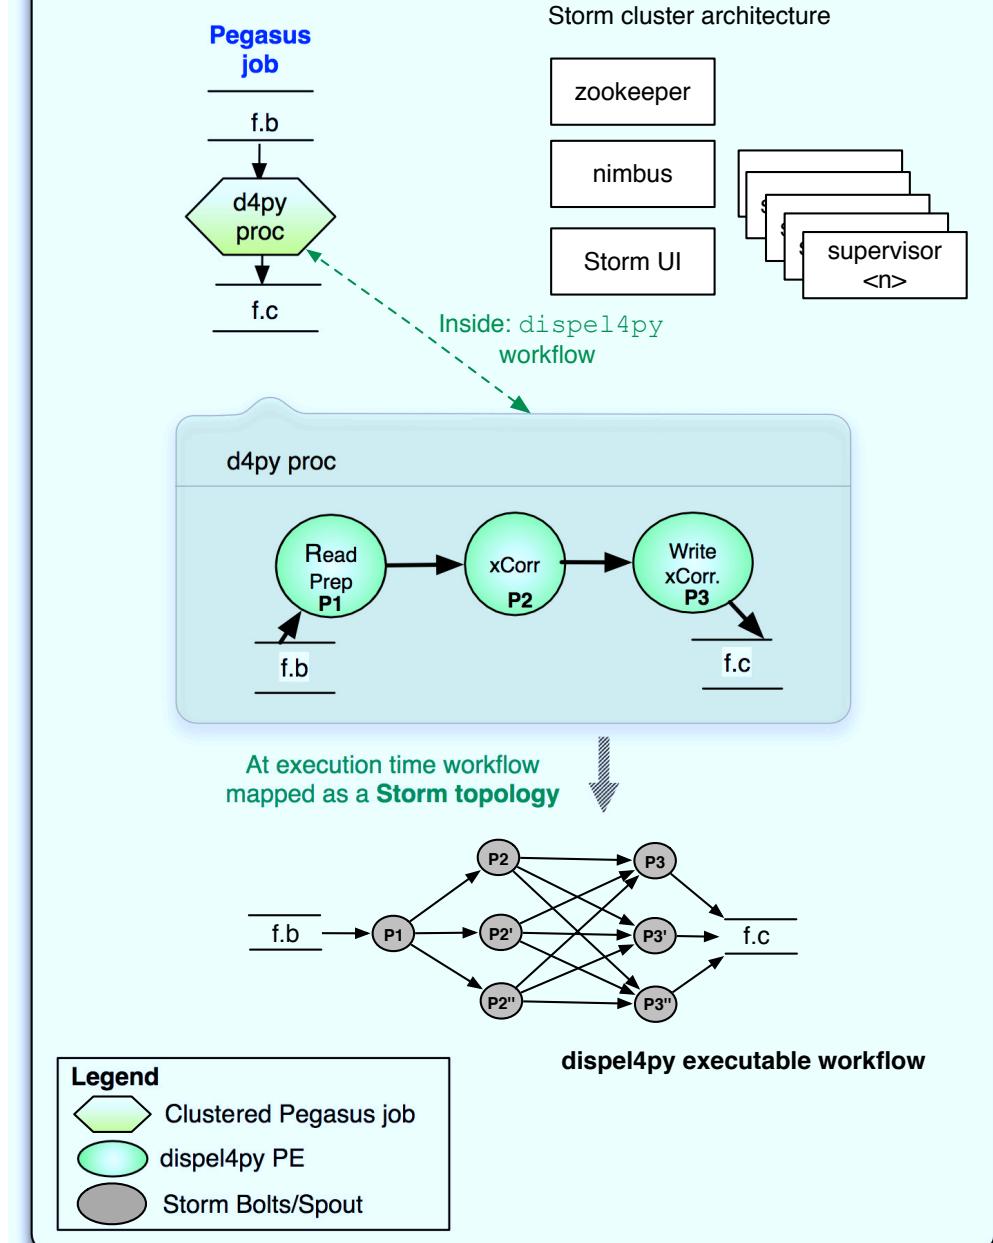
Container 1

Submit Asterism

3 instances as Container 3 (zookeeper, nimbus, Storm UI)

16 instances as Container 3 (Supervisors)

Execution environment -- Container 3- Storm, dispel4py, Obspy



NSF-Chameleon cloud
1 node – 40 cores



Asterism Evaluations

Experiment 1: Data from IRIS services (394 stations)

Time

Phase 1 – 8 minutes

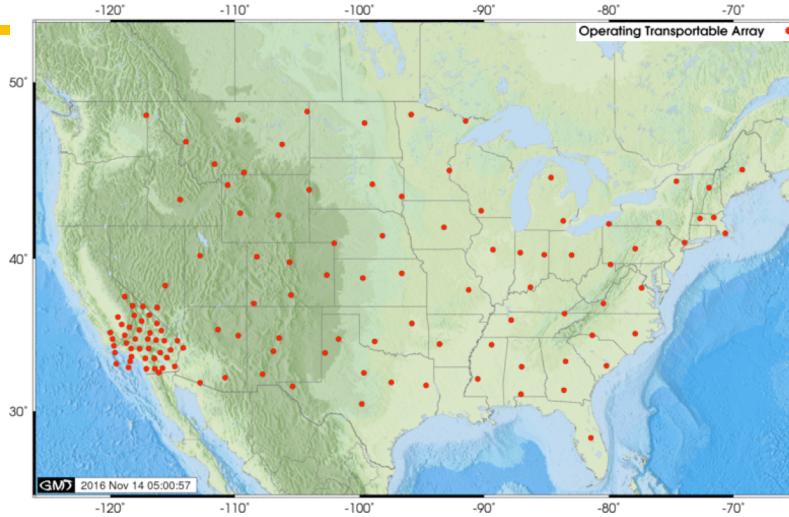
Phase 2 – 2 hours

Moving data < 1 minute

Data size

Input data 150MB

Output data 39GB



Experiment 2: Workflow for 3 days requesting data every 2 hours

Scope of this work

Executing & paralyzing automatically data-intensive applications

in heterogeneous systems with different enactment engines

NSF-
Chameleon
cloud
1 node –
40 cores

Maximizing performance

- both phases in the MPI cluster
- increasing the number of Storm Supervisors

dispel4py Performance Evaluations

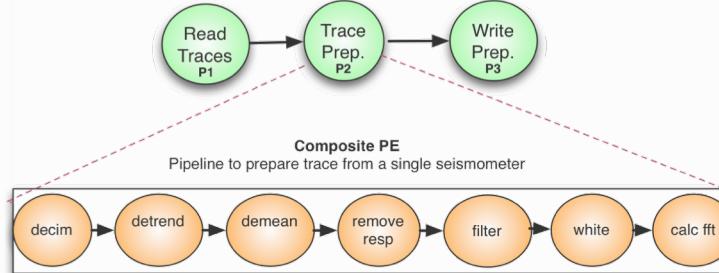


Computing Resources	Terracorrelator	SuperMuc	Amazon EC2	EDIM1
Type	Shared-memory	Cluster	Cloud	Cloud
Enactment Systems	MPI, multi	MPI, multi	MPI, Storm, multi	MPI, Storm, multi
Nodes	1	16	18	14
Cores per Node	32	16	2	4
Total Cores	32	256	36	14
Memory	2TB	32GB	4GB	3GB

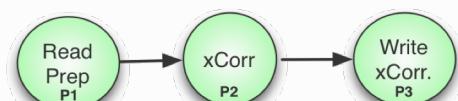
dispel4py Performance Evaluations



Data-Intensive Application



dispel4py preproc. (Phase 1)



dispel4py proc. (Phase 2)

1000 stations
Input 150MB
Output 39GB

Both workflows (Phase 1 and Phase 2) executed in the same computing resource and the same mapping.

Mode	Terracorrelator (32 cores/Node)	SuperMuc (256 cores 16 cores/node)	Amazon (36 cores 2 cores/node)	EDIM1 (14 cores 4 cores/node)
MPI	1501.32 (~25minutes)	1093.16 (~19minutes)	16862.73 (~5hours)	38656.94 (~11 hours)
multi	1332 .20 (~23minutes)			
Storm			27898.89 (~8 hours)	120077.123 (~33 hours)

Asterism: Easy-to-use system to empower data-driven research

New framework that automatically

manage the entire workflow, monitor its execution

handle data transfers between different platforms

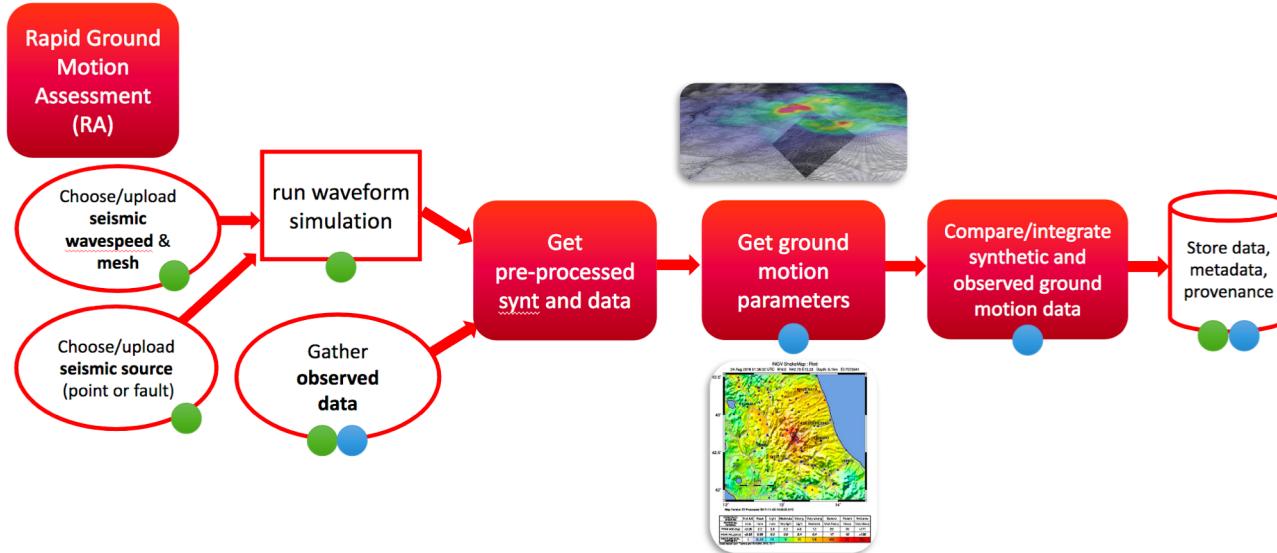
map to different enactment engines at runtime

DARE (Delivering Agile Research Excellence on European e-Infrastructures)

Provide scientific communities with tools/frameworks/APIs to allow for data-driven experiments, and rapid prototyping:

Address the requirements of European RIs, initially of EPOS, on Earth science, and IS/ENES2, on climate.

Rapid Ground Motion Assessment (RA)



Aims:

Model the strong ground motion after large earthquakes

Rapid assessment of earthquakes' impact, and emergency response

Steps:

Build dispel4py workflows to test them locally – using small dataset.

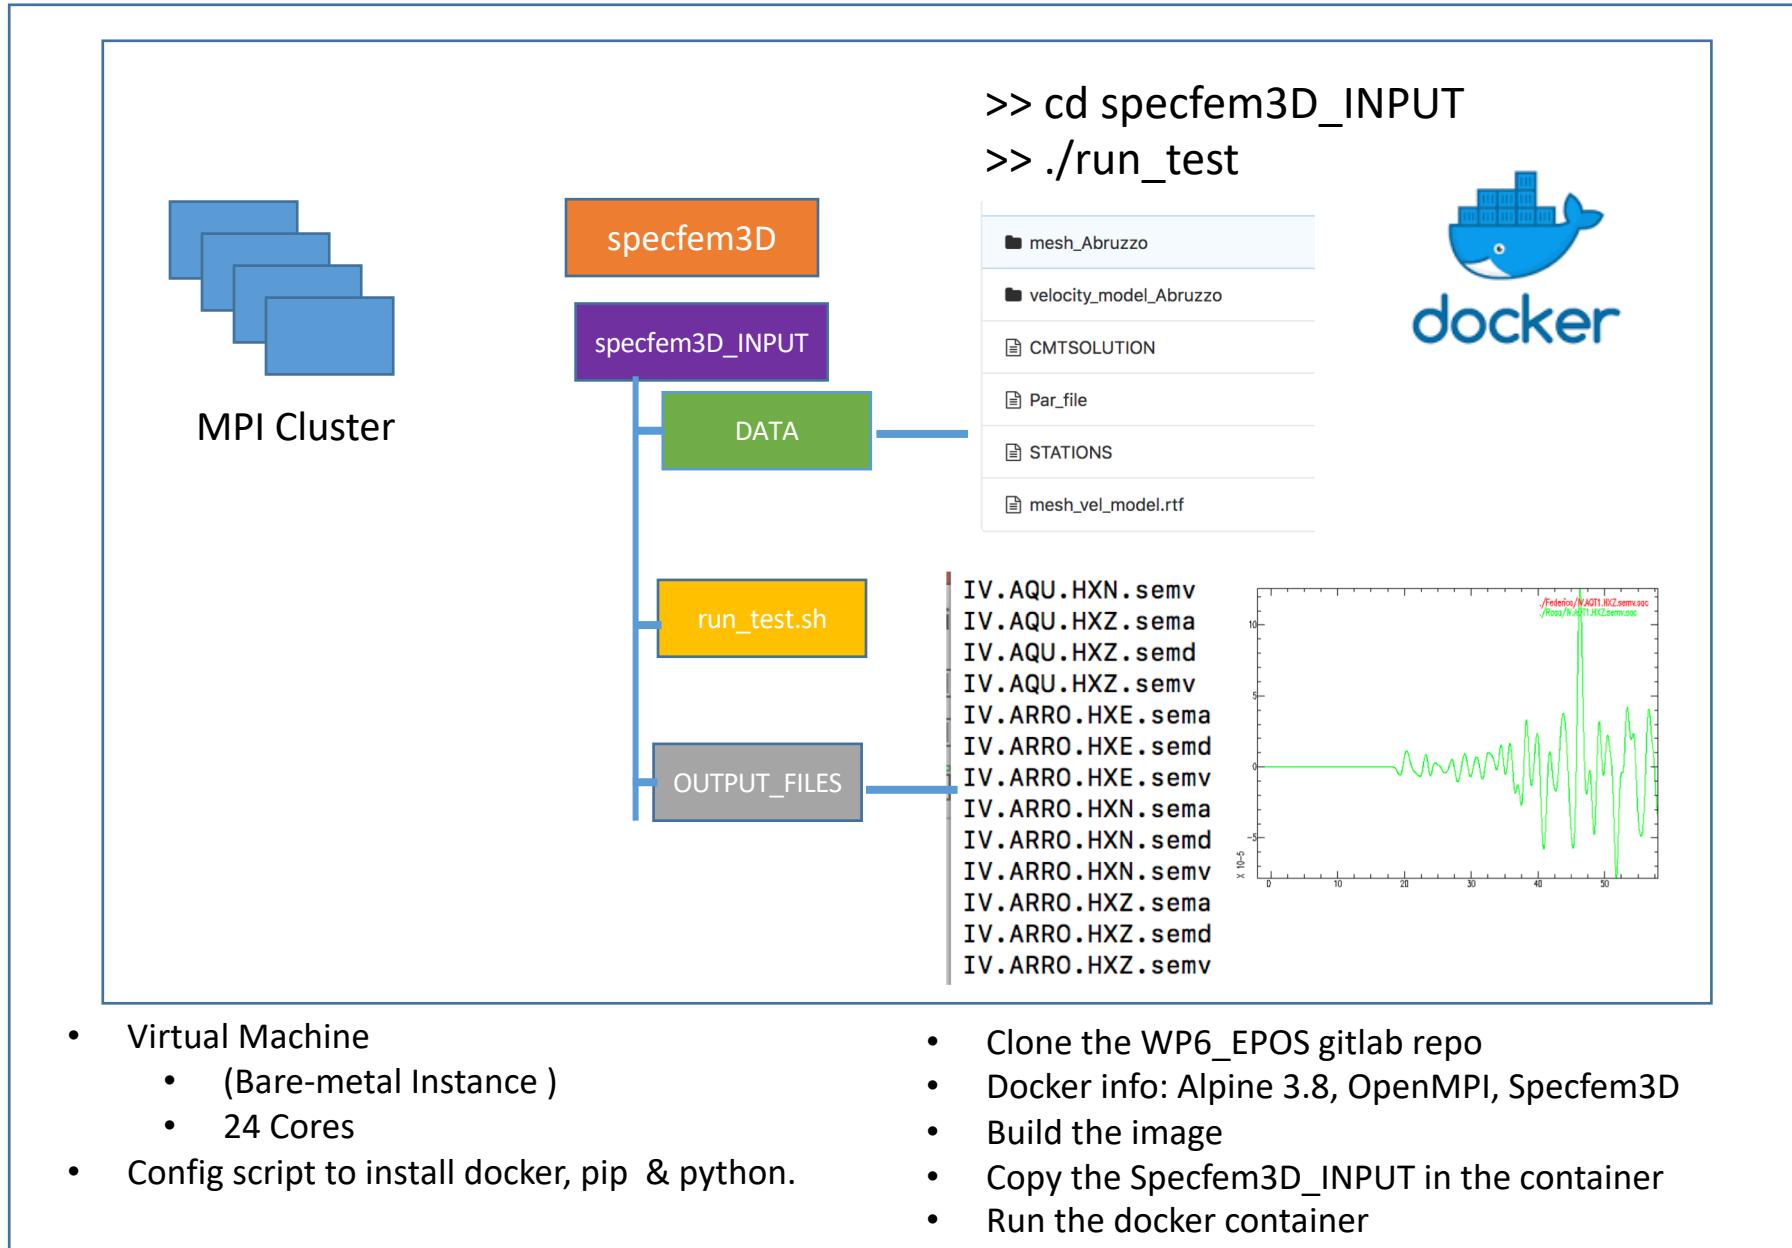
Scale them up:

- Run the workflows using HPC/Cloud
 - ** Using dispel4py parallel mappings

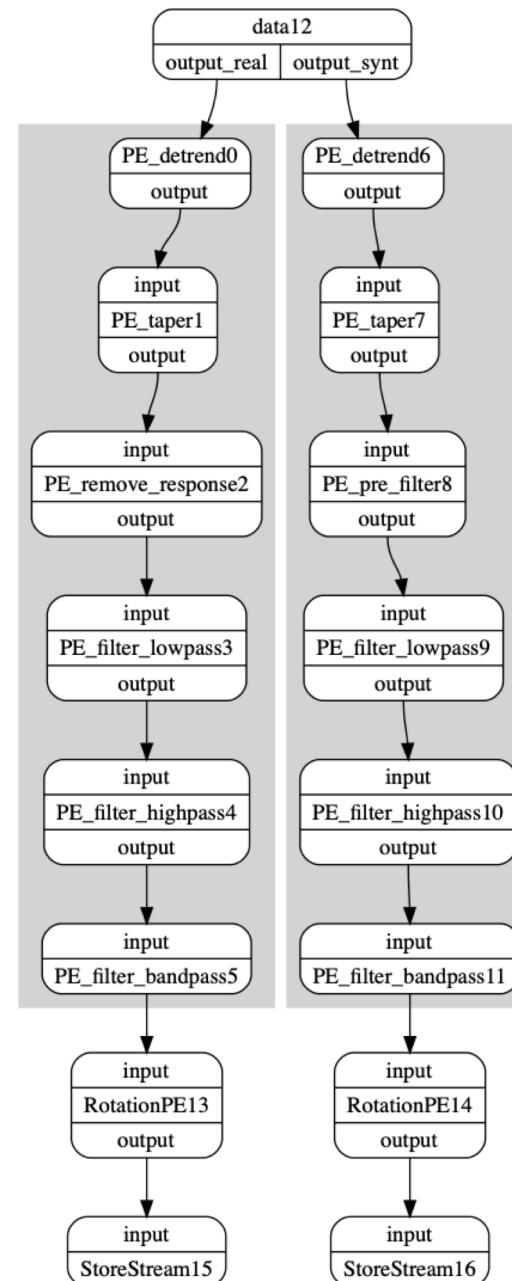
Waveform simulation: Specfem3D + MPI cluster

Test Case: RA

run waveform simulation



Get Pre-Preprocessed Synthetic and Data



Test Case: RA

Get
pre-processed
synt and data

Get Pre-Preprocessed Synthetic and Data



run_preprocess_misfit.sh 175 Bytes

```
1 #!/bin/bash
2
3 export PYTHONPATH=$PYTHONPATH
4 export MISFIT_PREP_CONFIG="pr
5 echo $MISFIT_PREP_CONFIG
6 dispel4py simple create_misfi
7
```

data

IV.ARRO..EHE.mseed
IV.ARRO..HNH.mseed
IV.ARRO..HZH.mseed

synth

IV.ARRO.HXE.sema
IV.ARRO.HXE.semdu
IV.ARRO.HXE.semva
IV.ARRO.HXN.sema
IV.ARRO.HXN.semdu
IV.ARRO.HXN.semva
IV.ARRO.HXZ.sema
IV.ARRO.HXZ.semdu
IV.ARRO.HXZ.semva

output

misfit_input.json 848 Bytes

```
1 {
2     "data": [
3         {
4             "input": {
5                 "data_dir": "/Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/misfit_data/data",
6                 "synt_dir": "/Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/misfit_data/synth",
7                 "events": "/Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/misfit_data/events_simulation_CI_CI_test_0_1507128030823",
8                 "event_id": "smi:webservices.ingv.it/fdsnws/event/1/query?eventId=1744261",
9                 "stations_dir": "/Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/misfit_data/stations",
10                "output_dir" : "/Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/misfit_data/output",
11                "network": [
12                    "IV"
13                ],
14                "station": [
15                    "ARRO"
16                ]
17            }
18        }
19    ]
20}
```

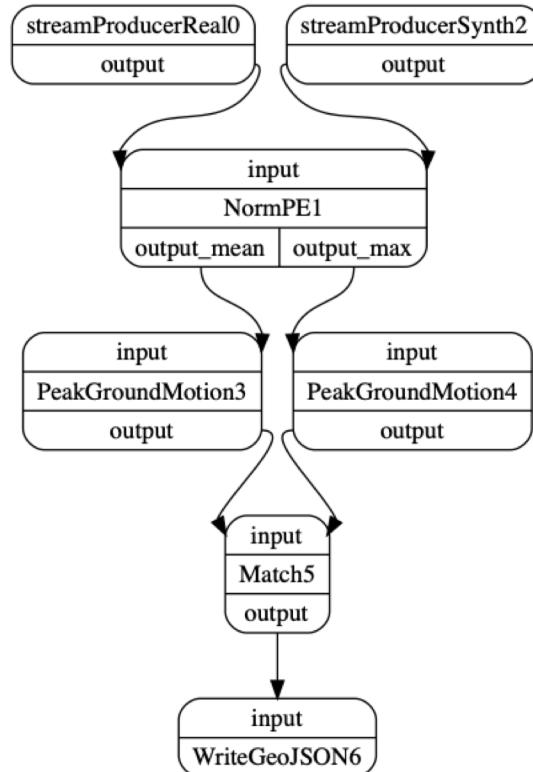
IV.ARRO.HXZ.synth
IV.ARRO.HXR.synth
IV.ARRO.HXT.synth
IV.ARRO.EHZ.data
IV.ARRO.EHR.data
IV.ARRO.EHT.data

Pre-preprocessed synthetic and observed data
(Underline files are used in the next step (slide))

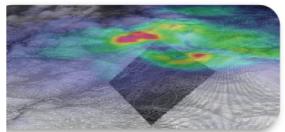
Note: Those are the synth seismograms generated in the previous step (slide)



Get and Compare Ground Motion Parameters



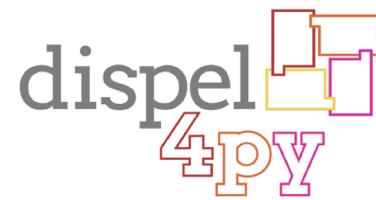
Test Case: RA



Get ground motion parameters

Compare/integrate synthetic and observed ground motion data

Get and Compare Ground Motion Parameters



run_RA.sh 210 Bytes

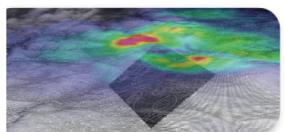
Lock Edit Web IDE Replace Delete

```
1 dispel4py simple dispel4py_RA.pgm_story.py -d '{"streamProducerReal": [ {"input": "./misfit_data/output/IV.ARRO.EHR.data"} ], "streamProducerSynth":
```

output

ARRO_max.json ARRO_mean.json

Test Case: RA



Get ground motion parameters

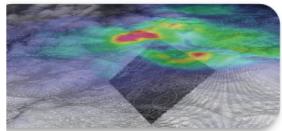
Compare/integrate synthetic and observed ground motion data

Get and Compare Ground Motion Parameters



<pre>type: "Feature" properties: station: "ARRO" data: PGD: 0.00003058970011915515 PGV: 0.00015387362683992627 PGA: 0.0008035731527687157 p_norm: "max" PSA_0.3Hz: 0.000171880777011629 PSA_1.0Hz: 0.0020140831084336057 PSA_3.0Hz: 0.0009812436987400703 synt: PGD: 0.000038899272805005096 PGV: 0.00011736045694475592 PGA: 0.00035493665398226155 p_norm: "max" PSA_0.3Hz: 0.0003942182089108563 PSA_1.0Hz: 0.0007176179285863899 PSA_3.0Hz: 0.0003707452407479844 difference: PGD: -0.000008309572685849946 PGV: 0.000036513169895170355 PGA: 0.0004486364987864541 PSA_0.3Hz: -0.000223374318992273 PSA_1.0Hz: 0.0012964651798472158 PSA_3.0Hz: 0.0006104984579920859 relative_difference: PGD: -0.2716460983102781 PGV: 0.23729322980834633 PGA: 0.5583020005592205 PSA_0.3Hz: -1.2935561251517065 PSA_1.0Hz: 0.6436999418834826 PSA_3.0Hz: 0.6221680289778919 geometry: type: "Point" coordinates: [] </pre>	<pre>type: "Feature" properties: station: "ARRO" data: PGD: 0.00003058970011915515 PGV: 0.00015387362683992627 PGA: 0.0008035731527687157 p_norm: "mean" PSA_0.3Hz: 0.000171880777011629 PSA_1.0Hz: 0.0020140831084336057 PSA_3.0Hz: 0.0009812436987400703 synt: PGD: 0.000038899272805005096 PGV: 0.00011736045694475592 PGA: 0.00035493665398226155 p_norm: "mean" PSA_0.3Hz: 0.0003942182089108563 PSA_1.0Hz: 0.0007176179285863899 PSA_3.0Hz: 0.0003707452407479844 difference: PGD: -0.000008309572685849946 PGV: 0.000036513169895170355 PGA: 0.0004486364987864541 PSA_0.3Hz: -0.000223374318992273 PSA_1.0Hz: 0.0012964651798472158 PSA_3.0Hz: 0.0006104984579920859 relative_difference: PGD: -0.2716460983102781 PGV: 0.23729322980834633 PGA: 0.5583020005592205 PSA_0.3Hz: -1.2935561251517065 PSA_1.0Hz: 0.6436999418834826 PSA_3.0Hz: 0.6221680289778919 geometry: type: "Point" coordinates: [] </pre>
---	---

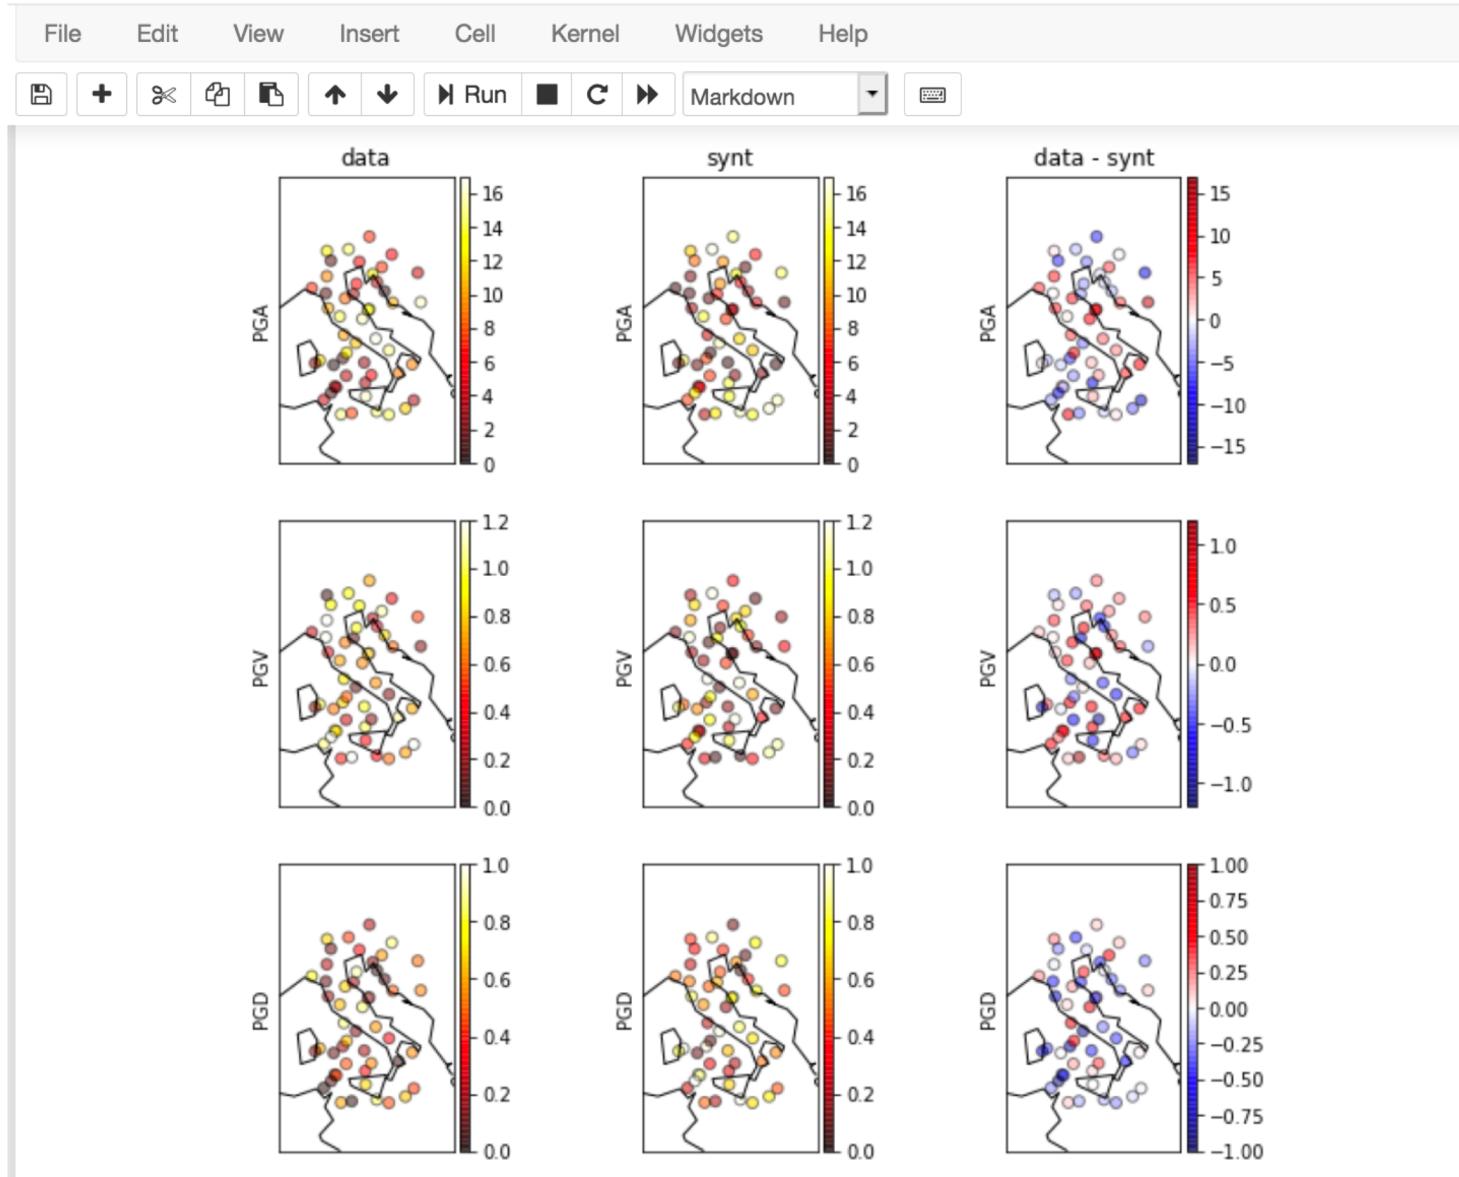
Test Case: RA



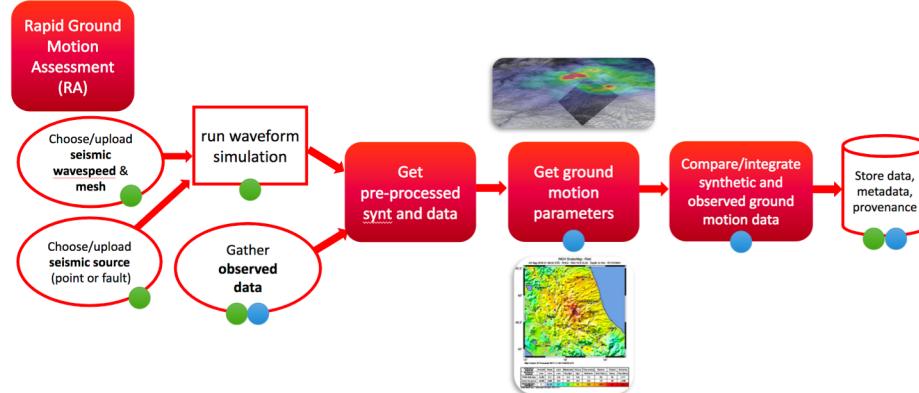
Get ground motion parameters

Compare/integrate synthetic and observed ground motion data

RA Maps



Rapid Ground Motion Assessment (RA)



dispel4py +

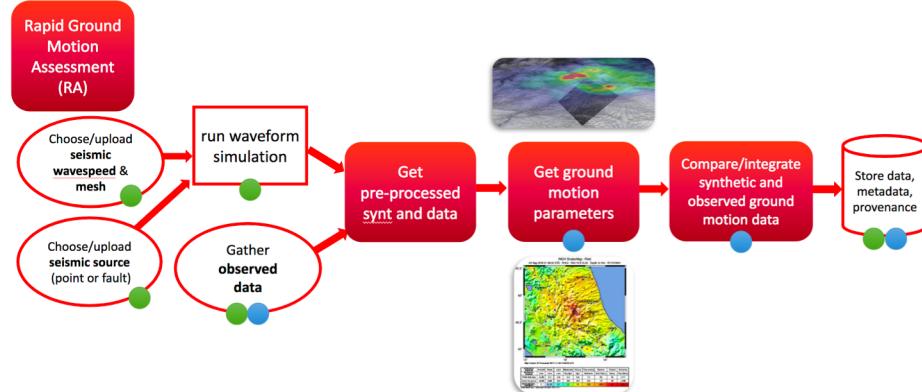


→ semantics and description

CWL is a specification for describing the data and execution model of workflows/command tools

Provides a consistent way to connect programs

Rapid Ground Motion Assessment (RA)



dispel4py +



→ semantics and description

CWL is a specification for describing the data and execution model of workflows/command tools

Provides a consistent way to connect programs

dispel4py-RA-pgm_story-job.yml 372 Bytes

```

1 script:
2   class: File
3   path: /Users/rosafilgueira/EPCC/DARE/WP6/test/MISFIT_RA/dispel4py_RA.pgm_story.py
4
5
6 input: '{"streamProducerReal": [ {"input": "/Users/rosafilgueira/EPCC/DARE/WP6/test/MI
  
```

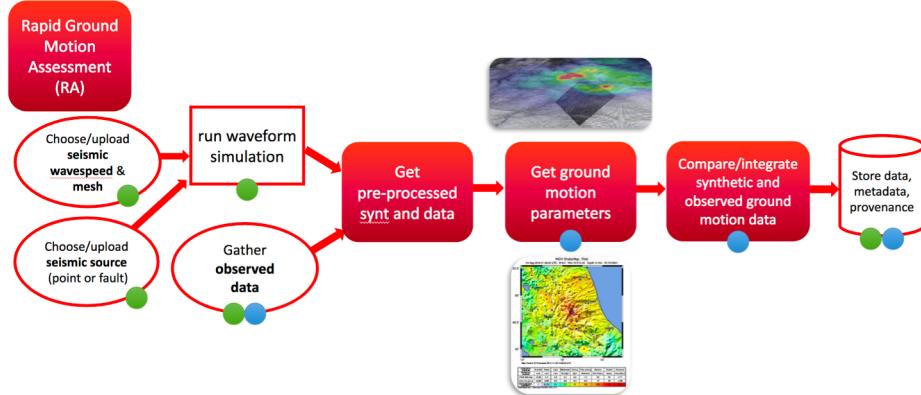
dispel4py-RA-pgm_story.cwl 437 Bytes

```

1 #!/usr/bin/env cwl-runner
2
3 cwlVersion: v1.0
4 class: CommandLineTool
5 baseCommand: [dispel4py, simple]
6 requirements:
7   EnvVarRequirement:
8     envDef:
9       PYTHONPATH: ${inputs.script.dirname}
10
11 inputs:
12   - id: script
13     type: File
14     inputBinding:
15       position: 1
16   - id: input
17     type: string
18     inputBinding:
19       prefix: -d
20       position: 2
21
22 outputs:
23   output:
24     type:
25       type: array
26       items: File
27     outputBinding:
28       glob: "*.json"
  
```

cwl-runner dispel4py-RA-pgm_story.cwl dispel4py-RA-pgm_story-job.yml

Rapid Ground Motion Assessment (RA)



Once tested all the workflows locally (seq. map),
They can be scaled up automatically → Next step.

Provenance data is also collected in runtime (optional)

Metadata, Input and Output files --> registry/catalog

Future development:

Optimization of workflows → dynamic deployment
→ mapping selection

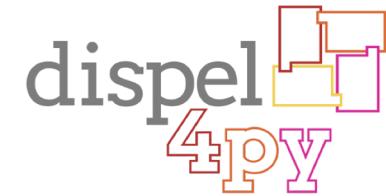
Delivering easy-to-use frameworks to empower data-driven research

Dr. Rosa Filgueira, University of Edinburgh, EPCC



Extra Slides

dispel4py parallel stream-based dataflow system



```
from dispel4py.workflow_graph import WorkflowGraph

pe1 = filterTweet ()
pe2 = counterHashTag ()
pe3 = counterLanguage ()
pe4 = statistics ()

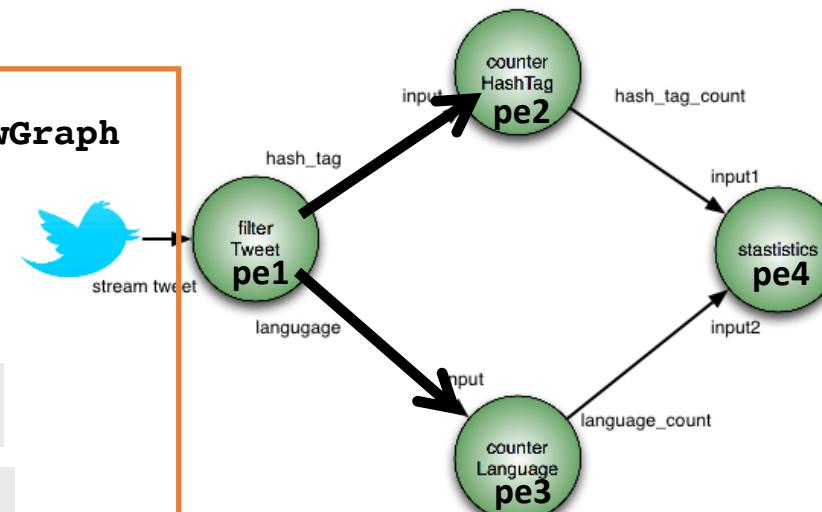
graph = WorkflowGraph ( )

graph.connect(pe1,'hash_tag',pe2,'input') ←
graph.connect(pe1,'language',pe3,'input') ←
graph.connect(pe2,'hash_tag_count',pe4,'input1')
graph.connect(pe3,'language_count',pe4,'input2')
```

PEs objects

Graph

Connections



Users only have to implement:

- PEs
- Connections



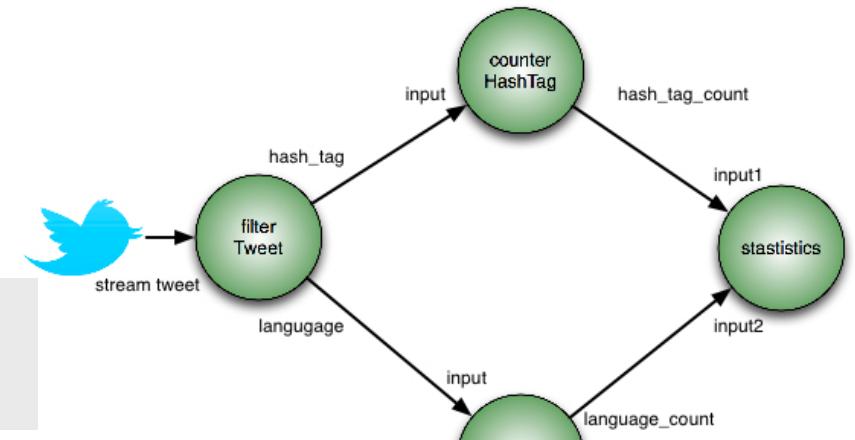
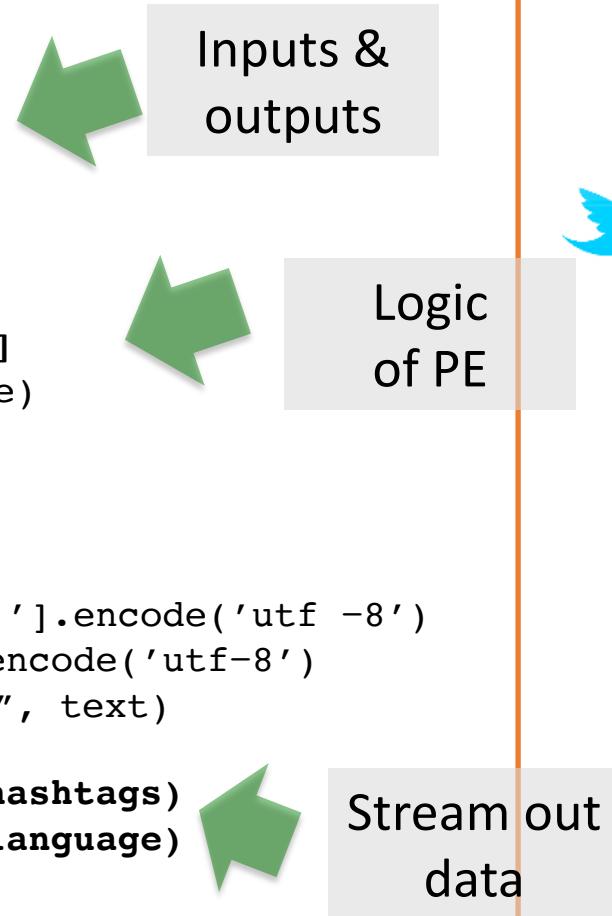
|epcc

dispelp4y basic concepts– Example of a PE

```
Class filterTweet(GenericPE):
def __init__( self ):
    GenericPE.init (self)
    self.add_output('hash_tags ')
    self.add_output('language')

def process ( self , inputs ):
    twitterDataFile= inputs['input ']
    tweet file = open(twitterDataFile)
    for line in tweet file :
        tweet = json.loads( line )
        language = ' '
        hashtags =[]
        language = tweet[u'lang '].encode('utf -8')
        text = tweet[u'text '].encode('utf-8')
        hashtags=re.findall(r"#(\w+)", text)

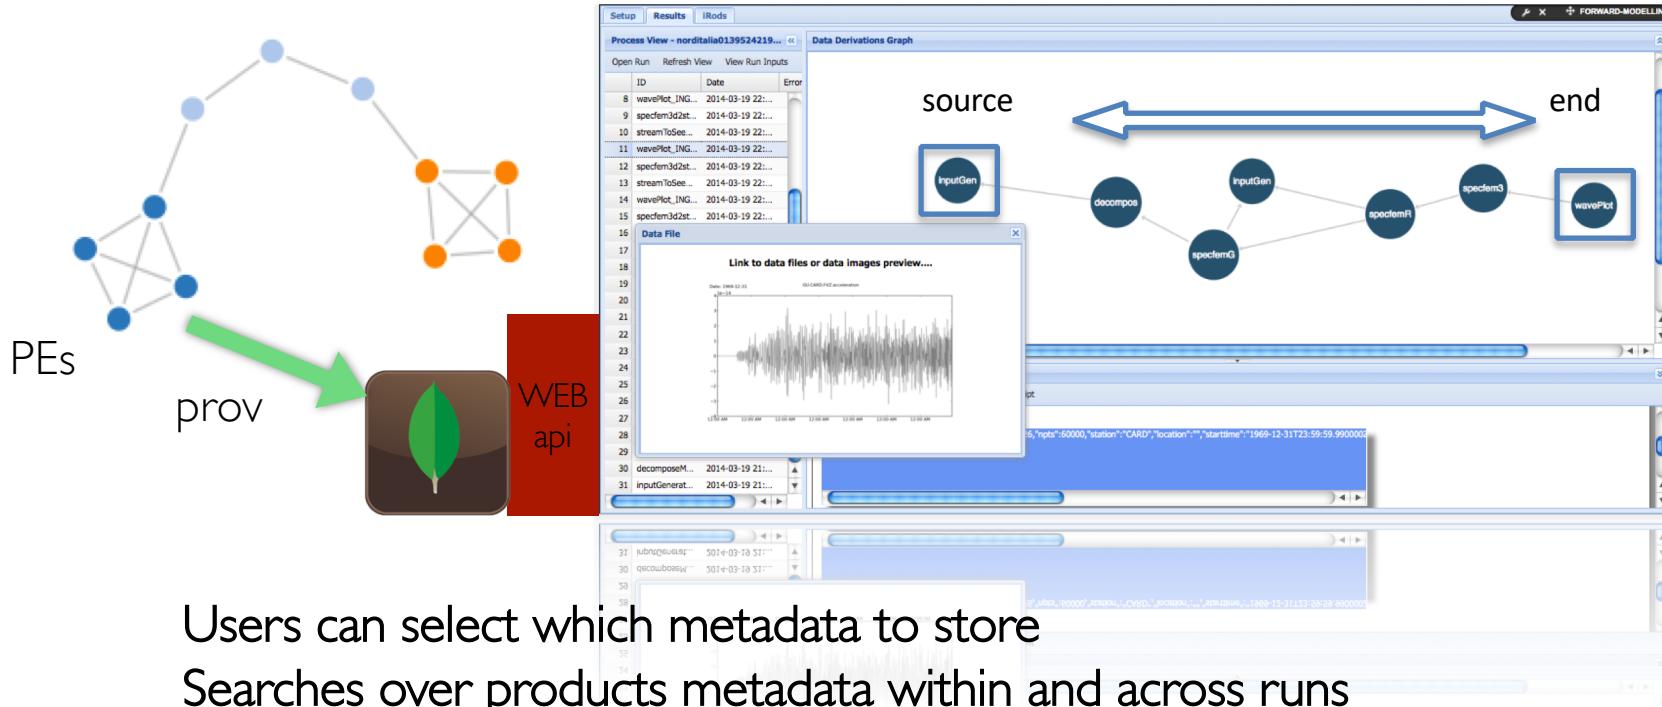
        self.write('hash tag', hashtags)
        self.write('language', language)
```



Users only have to implement:

- PEs
- Connections

dispel4py advance concepts – Provenance



Users can select which metadata to store
Searches over products metadata within and across runs
Data download and preview
Capturing of Errors for Diagnostic purposes
Data Fabric: Multi directional navigations across data dependencies
W3C PROV-DM as reference model.